


CSI 62 Operating Systems and Systems Programming Lecture 7

Synchronization

February 10, 2016
Prof. David Culler
<http://cs162.eecs.Berkeley.edu>

Recall: How does Thread get started?



- Eventually, `run_new_thread()` will select this TCB and return into beginning of `ThreadRoot()`
 - This really starts the new thread

2/10/16

Joseph CSI 62 ©UCB Spring 2016

Lec 7.2

Goals for Today

- Synchronization Operations
- Higher-level Synchronization Abstractions
 - Semaphores, monitors, and condition variables
- Programming paradigms for concurrent programs




2/10/16

Joseph CSI 62 ©UCB Spring 2016

Lec 7.3

Recall: Thread Abstraction



- Infinite number of processors
- Threads execute with variable speed
 - Programs must be designed to work with any schedule


2/10/16

Joseph CSI 62 ©UCB Spring 2016

Lec 7.4

Recall: Multiprocessing vs Multiprogramming

- What does it mean to run two threads "concurrently"?
 - Scheduler is free to run threads in any order and interleaving: FIFO, Random, ...
 - Dispatcher can choose to run each thread to completion or time-slice in big chunks or small chunks



- Also recall: Hyperthreading
 - Possible to interleave threads on a per-instruction basis
 - Keep this in mind for our examples (like multiprocessing)

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.5

Correctness for systems with concurrent threads


- If dispatcher can schedule threads in any way, programs must work under all circumstances
 - Can you test for this?
 - How can you know if your program works?
- Independent Threads:**
 - No state shared with other threads
 - Deterministic \Rightarrow Input state determines results
 - Reproducible \Rightarrow Can recreate Starting Conditions, I/O
 - Scheduling order doesn't matter (if `switch()` works!!!)
- Cooperating Threads:**
 - Shared State between multiple threads
 - Non-deterministic
 - Non-reproducible
- Non-deterministic and Non-reproducible means that bugs can be intermittent
 - Sometimes called "Heisenbugs"

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.6

High-level Example: Web Server



- Server must handle many requests
- Non-cooperating version:


```
serverLoop() {
    con = AcceptCon();
    ProcessWebPage(con);
}
```
- What are some disadvantages of this technique?

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.7

Threaded Web Server

- Now, use a single process
- Multithreaded (cooperating) version:


```
serverLoop() {
    connection = AcceptCon();
    ThreadFork(ServiceWebPage(), connection);
}
```
- Looks almost the same, but has many advantages:
 - Can share file caches kept in memory, results of CGI scripts, other things
 - Threads are *much* cheaper to create than processes, so this has a lower per-request overhead
- Question: would a user-level (say one-to-many) thread package make sense here?
 - When one request blocks on disk, all block...
- What about Denial of Service attacks or digg / Slash-dot effects?

2/10/16


Joseph CS162 ©UCB Spring 2016



Lec 7.8

Thread Pools

- Problem with previous version: Unbounded Threads
 - When web-site becomes too popular – throughput sinks
- Instead, allocate a bounded “pool” of worker threads, representing the maximum level of multiprogramming



```

master() {
    allocThreads(worker, queue);
    while (TRUE) {
        con=AcceptCon();
        Enqueue(queue, con);
        wakeUp(queue);
    }
}


worker(queue) {
    while (TRUE) {
        con=Dequeue(queue);
        if (con==null)
            sleepOn(queue);
        else
            ServiceWebPage(con);
    }
}
    
```

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.9

ATM Bank Server



- ATM server problem:
 - Service a set of requests
 - Do so without corrupting database
 - Don't hand out too much money

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.10

ATM bank server example

- Suppose we wanted to implement a server process to handle requests from an ATM network:

```

BankServer() {
    while (TRUE) {
        ReceiveRequest(&op, &acctId, &amount);
        ProcessRequest(op, acctId, amount);
    }
}

ProcessRequest(op, acctId, amount) {
    if (op == deposit) Deposit(acctId, amount);
    else if ...
}

Deposit(acctId, amount) {
    acct = GetAccount(acctId); /* may use disk I/O */
    acct->balance += amount;
    StoreAccount(acct); /* Involves disk I/O */
}
    
```

- How could we speed this up?
 - More than one request being processed at once
 - Event driven (overlap computation and I/O)
 - Multiple threads (multi-proc, or overlap comp and I/O)

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.11

Event Driven Version of ATM server

- Suppose we only had one CPU
 - Still like to overlap I/O with computation
 - Without threads, we would have to rewrite in event-driven style
- Example


```

BankServer() {
    while (TRUE) {
        event = WaitForNextEvent();
        if (event == ATMRequest)
            StartOnRequest();
        else if (event == AcctAvail)
            ContinueRequest();
        else if (event == AcctStored)
            FinishRequest();
    }
}
            
```

 - What if we missed a blocking I/O step?
 - What if we have to split code into hundreds of pieces which could be blocking?
 - This technique is used for graphical programming

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.12

Can Threads Make This Easier?

- Threads yield overlapped I/O and computation without “deconstructing” code into non-blocking fragments
 - One thread per request
- Requests proceeds to completion, blocking as required:

```
Deposit(acctId, amount) {  
    acct = GetAccount(acctId); /* May use disk I/O */  
    acct->balance += amount;  
    StoreAccount(acct);      /* Involves disk I/O */  
}
```

- Unfortunately, shared state can get corrupted:

<u>Thread 1</u>	<u>Thread 2</u>
load r1, acct->balance	
	load r1, acct->balance
	add r1, amount2
	store r1, acct->balance
add r1, amount1	
store r1, acct->balance	

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.13

Problem is at the lowest level

- Most of the time, threads are working on separate data, so scheduling doesn't matter:

Thread A
x = 1;

Thread B
y = 2;

- However, What about (Initially, y = 12):

Thread A
x = 1;
x = y+1;

Thread B
y = 2;
y = y*2;

- What are the possible values of x?
- Or, what are the possible values of x below?

Thread A
x = 1;

Thread B
x = 2;

- X could be 1 or 2 (non-deterministic)
- Could even be 3 for serial processors:
 - Thread A writes 0001, B writes 0010.
 - Scheduling order ABABABBA yields 3!

♪

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.14

Atomic Operations

- To understand a concurrent program, we need to know what the underlying indivisible operations are!
- Atomic Operation:** an operation that always runs to completion or not at all
 - It is *indivisible*: it cannot be stopped in the middle and state cannot be modified by someone else in the middle
 - Fundamental building block – if no atomic operations, then have no way for threads to work together
- On most machines, memory references and assignments (i.e. loads and stores) of words are atomic
 - Consequently – weird example that produces “3” on previous slide can't happen
- Many instructions are not atomic
 - Double-precision floating point store often not atomic
 - VAX and IBM 360 had an instruction to copy a whole array

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.15

Correctness Requirements

- Threaded programs must work for all interleavings of thread instruction sequences
 - Cooperating threads inherently non-deterministic and non-reproducible
 - Really hard to debug unless carefully designed!

- Example: Therac-25

- Machine for radiation therapy
 - Software control of electron accelerator and electron beam/Xray production
 - Software control of dosage
- Software errors caused the death of several patients
 - A series of race conditions on shared variables and poor software design




Figure 1. Typical Therac-25 facility

- “They determined that data entry speed during editing was the key factor in producing the error condition: If the prescription data was edited at a fast pace, the overdose occurred.”


2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.16

Space Shuttle Example

- Original Space Shuttle launch aborted 20 minutes before scheduled launch
- Shuttle has five computers:
 - Four run the "Primary Avionics Software System" (PASS)
 - Asynchronous and real-time
 - Runs all of the control systems
 - Results synchronized and compared every 3 to 4 ms
 - The Fifth computer is the "Backup Flight System" (BFS)
 - stays synchronized in case it is needed
 - Written by completely different team than PASS
- Countdown aborted because BFS disagreed with PASS
 - A 1/67 chance that PASS was out of sync one cycle
 - Bug due to modifications in **initialization** code of PASS
 - A delayed init request placed into timer queue
 - As a result, timer queue not empty at expected time to force use of hardware clock
 - Bug not found during extensive simulation



2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.17

Another Concurrent Program Example

- Two threads, A and B, compete with each other
 - One tries to increment a shared counter
 - The other tries to decrement the counter

Thread A

```
i = 0;
while (i < 10)
  i = i + 1;
printf("A wins!");
```

Thread B

```
i = 0;
while (i > -10)
  i = i - 1;
printf("B wins!");
```

- Assume that memory loads and stores are atomic, but incrementing and decrementing are **not** atomic
- Who wins? Could be either
- Is it guaranteed that someone wins? Why or why not?
- What if both threads have their own CPU running at same speed? Is it guaranteed that it goes on forever?

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.18

Hand Simulation Multiprocessor Example

- Inner loop looks like this:

<u>Thread A</u>		<u>Thread B</u>	
r1=0	load r1, M[i]	r1=0	load r1, M[i]
r1=1	add r1, r1, 1	r1=-1	sub r1, r1, 1
M[i]=1	store r1, M[i]	M[i]=-1	store r1, M[i]

- Hand Simulation:**
 - And we're off. A gets off to an early start
 - B says "hmpf, better go fast" and tries really hard
 - A goes ahead and writes "1"
 - B goes and writes "-1"
 - A says "HUH??? I could have sworn I put a 1 there"
- Could this happen on a uniprocessor?
 - Yes! Unlikely, but if you are depending on it not happening, it will and your system will break...

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.19

Administrivia

- Group/Section assignments should be completed!
 - If you are not in group, talk to us immediately!
- Section assignments out on piazza
 - Start going to them this week
 - Need to know your TA!
 - Participation is 5% of your grade
 - Should attend section with your TA
- First design doc due this Friday
 - This means you should be well on your way with Project 1
 - Watch for notification from your TA to sign up for design review

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.20

Motivation: "Too Much Milk"

- Great thing about OS's – analogy between problems in OS and problems in real life
 - Help you understand real life problems better
 - But, computers are much stupider than people
- Example: People need to coordinate:



Time	Person A	Person B
3:00	Look in Fridge. Out of milk	
3:05	Leave for store	
3:10	Arrive at store	Look in Fridge. Out of milk
3:15	Buy milk	Leave for store
3:20	Arrive home, put milk away	Arrive at store
3:25		Buy milk
3:30		Arrive home, put milk away

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.2.1

Definitions

- **Synchronization**: using atomic operations to ensure cooperation between threads
 - For now, only loads and stores are atomic
 - We are going to show that its hard to build anything useful with only reads and writes
- **Mutual Exclusion**: ensuring that only one thread does a particular thing at a time
 - One thread *excludes* the other while doing its task
- **Critical Section**: piece of code that only one thread can execute at once. Only one thread at a time will get into this section of code
 - Critical section is the result of mutual exclusion
 - Critical section and mutual exclusion are two ways of describing the same thing


2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.2.2

More Definitions

- **Lock**: prevents someone from doing something
 - Lock before entering critical section and before accessing shared data
 - Unlock when leaving, after accessing shared data
 - Wait if locked
- » Important idea: all synchronization involves waiting
- For example: fix the milk problem by putting a key on the refrigerator
 - Lock it and take key if you are going to go buy milk
 - Fixes too much: roommate angry if only wants OJ



- Of Course – We don't know how to make a lock yet

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.2.3

Too Much Milk: Correctness Properties

- Need to be careful about correctness of concurrent programs, since non-deterministic
 - Always write down behavior first
 - Impulse is to start coding first, then when it doesn't work, pull hair out
 - Instead, think first, then code
- What are the correctness properties for the "Too much milk" problem???
- Never more than one person buys
- Someone buys if needed
- Restrict ourselves to use only atomic load and store operations as building blocks

2/10/16


Joseph CS162 ©UCB Spring 2016

Lec 7.2.4

Too Much Milk: Solution #1

- Use a note to avoid buying too much milk:
 - Leave a note before buying (kind of “lock”)
 - Remove note after buying (kind of “unlock”)
 - Don't buy if note (wait)
- Suppose a computer tries this (remember, only memory read/write are atomic):

```
if (noMilk) {  
  if (noNote) {  
    leave Note;  
    buy milk;  
    remove note;  
  }  
}
```



- Result?
 - Still too much milk **but only occasionally!**
 - Thread can get context switched after checking milk and note but before buying milk!
- Solution makes problem worse since fails **intermittently**
 - Makes it really hard to debug...
 - Must work despite what the dispatcher does!

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.25

Too Much Milk: Solution #1 1/2

- Clearly the Note is not quite blocking enough
 - Let's try to fix this by placing note first
- Another try at previous solution:

```
leave Note;  
if (noMilk) {  
  if (noNote) {  
    leave Note;  
    buy milk;  
  }  
}  
remove note;
```

- What happens here?
 - Well, with human, probably nothing bad
 - With computer: no one ever buys milk



2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.26

Too Much Milk Solution #2

- How about labeled notes?
 - Now we can leave note before checking
- Algorithm looks like this:

```
Thread A  
leave note A;  
if (noNote B) {  
  if (noMilk) {  
    buy Milk;  
  }  
}  
remove note A;
```

```
Thread B  
leave note B;  
if (noNoteA) {  
  if (noMilk) {  
    buy Milk;  
  }  
}  
remove note B;
```


- Does this work?
- Possible for neither thread to buy milk
 - Context switches at exactly the wrong times can lead each to think that the other is going to buy
- Really insidious:
 - Extremely unlikely** that this would happen, but will at worse possible time
 - Probably something like this in UNIX

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.27

Too Much Milk Solution #2: problem!



- I'm not getting milk, You're getting milk
- This kind of lockup is called “starvation!”**

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.28

Too Much Milk Solution #3

- Here is a possible two-note solution:

Thread A	Thread B
leave note A;	leave note B;
while (note B) { //X	if (noNote A) { //Y
do nothing;	if (noMilk) {
}	buy milk;
if (noMilk) {	}
buy milk;	remove note B;
}	
remove note A;	

- Does this work? Yes. Both can guarantee that:
 - It is safe to buy, or
 - Other will buy, ok to quit
- At X:
 - if no note B, safe for A to buy,
 - otherwise wait to find out what will happen
- At Y:
 - if no note A, safe for B to buy
 - Otherwise, A is either buying or waiting for B to quit

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.29

Solution #3 discussion

- Our solution protects a single "Critical-Section" piece of code for each thread:

```
if (noMilk) {
    buy milk;
}
```

- Solution #3 works, but it's really unsatisfactory
 - Really complex – even for this simple an example
 - Hard to convince yourself that this really works
 - A's code is different from B's – what if lots of threads?
 - Code would have to be slightly different for each thread
 - While A is waiting, it is consuming CPU time
 - This is called "busy-waiting"
- There's a better way
 - Have hardware provide better (higher-level) primitives than atomic load and store
 - Build even higher-level programming abstractions on this new hardware support

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.30

Too Much Milk: Solution #4

- Suppose we have some sort of implementation of a lock (more in a moment).
 - Lock.Acquire()** – wait until lock is free, then grab
 - Lock.Release()** – Unlock, waking up anyone waiting
 - These must be atomic operations – if two threads are waiting for the lock and both see it's free, only one succeeds to grab the lock
- Then, our milk problem is easy:


```
milklock.Acquire();
if (nomilk)
    buy milk;
milklock.Release();
```
- Once again, section of code between **Acquire()** and **Release()** called a "**Critical Section**"
- Of course, you can make this even simpler: suppose you are out of ice cream instead of milk
 - Skip the test since you always need more ice cream.

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.31

Where are we going with synchronization?

Programs	Shared Programs			
Higher-level API	Locks	Semaphores	Monitors	Send/Receive
Hardware	Load/Store	Disable Ints	Test&Set	Comp&Swap

- We are going to implement various higher-level synchronization primitives using atomic operations
 - Everything is pretty painful if only atomic primitives are load and store
 - Need to provide primitives useful at user-level

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.32

How to implement Locks?

- **Lock**: prevents someone from doing something
 - Lock before entering critical section and before accessing shared data
 - Unlock when leaving, after accessing shared data
 - Wait if locked
 - » Important idea: all synchronization involves waiting
 - » Should sleep if waiting for a long time
- Atomic Load/Store: get solution like Milk #3
 - Looked at this last lecture
 - Pretty complex and error prone
- Hardware Lock instruction
 - Is this a good idea?
 - What about putting a task to sleep?
 - » How do you handle the interface between the hardware and scheduler?
 - Complexity?
 - » Done in the Intel 432 – each feature makes HW more complex and slow



2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.33

Naïve use of Interrupt Enable/Disable

- How can we build multi-instruction atomic operations?
 - Recall: dispatcher gets control in two ways.
 - » Internal: Thread does something to relinquish the CPU
 - » External: Interrupts cause dispatcher to take CPU
 - On a uniprocessor, can avoid context-switching by:
 - » Avoiding internal events (although virtual memory tricky)
 - » Preventing external events by disabling interrupts
- Consequently, naïve Implementation of locks:

```
LockAcquire { disable Ints; }
LockRelease { enable Ints; }
```
- Problems with this approach:
 - **Can't let user do this!** Consider following:

```
LockAcquire();
While(TRUE) {;
```
 - Real-Time system—no guarantees on timing!
 - » Critical Sections might be arbitrarily long
 - What happens with I/O or other important events?
 - » “Reactor about to meltdown. Help!”



2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.34

Better Implementation of Locks by Disabling Interrupts

- Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable

```
int value = FREE;
```



```
Acquire() {
    disable interrupts;
    if (value == BUSY) {
        put thread on wait queue;
        Go to sleep();
        // Enable interrupts?
    } else {
        value = BUSY;
    }
    enable interrupts;
}

Release() {
    disable interrupts;
    if (anyone on wait queue) {
        take thread off wait queue;
        Place on ready queue;
    } else {
        value = FREE;
    }
    enable interrupts;
}
```

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.35

New Lock Implementation: Discussion

- Why do we need to disable interrupts at all?
 - Avoid interruption between checking and setting lock value
 - Otherwise two threads could think that they both have lock

```
Acquire() {
    disable interrupts;
    if (value == BUSY) {
        put thread on wait queue;
        Go to sleep();
        // Enable interrupts?
    } else {
        value = BUSY;
    }
    enable interrupts;
}
```

} Critical Section

- Note: unlike previous solution, the critical section (inside **Acquire()**) is very short
 - User of lock can take as long as they like in their own critical section: doesn't impact global machine behavior
 - Critical interrupts taken in time!

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.36

Interrupt re-enable in going to sleep

- What about re-enabling ints when going to sleep?

```
Acquire() {  
    disable interrupts;  
    if (value == BUSY) {  
        put thread on wait queue;  
        Go to sleep();  
    } else {  
        value = BUSY;  
    }  
    enable interrupts;  
}
```

Enable Position
Enable Position
Enable Position

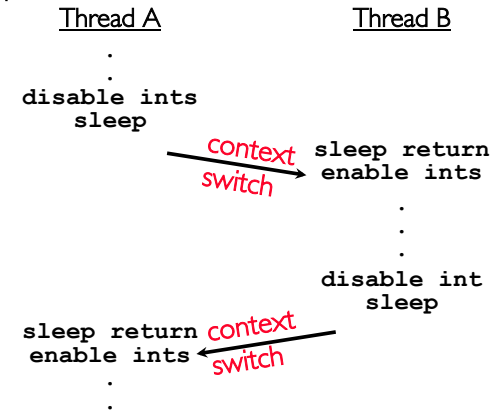
2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.37

How to Re-enable After Sleep()

- In scheduler, since interrupts are disabled when you call sleep:
 - Responsibility of the next thread to re-enable ints
 - When the sleeping thread wakes up, returns to acquire and re-enables interrupts



2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.38

Summary

- Important concept: Atomic Operations
 - An operation that runs to completion or not at all
 - These are the primitives on which to construct various synchronization primitives
- Talked about hardware atomicity primitives:
 - Disabling of Interrupts, test&set, swap, comp&swap, load-linked/store conditional
- Showed several constructions of Locks
 - Must be very careful not to waste/tie up machine resources
 - Shouldn't disable interrupts for long
 - Shouldn't spin wait for long
 - Key idea: Separate lock variable, use hardware mechanisms to protect modifications of that variable

2/10/16

Joseph CS162 ©UCB Spring 2016

Lec 7.39