

# CS162 Operating Systems and Systems Programming Lecture 9

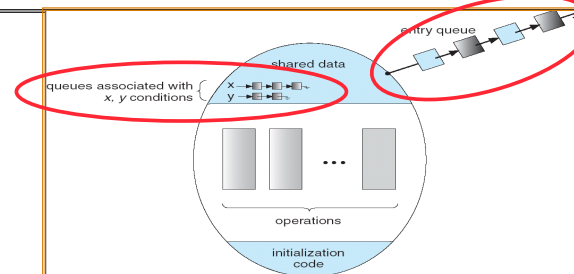
## Synchronization, Readers/Writers example, Scheduling

February 22<sup>nd</sup>, 2016

Prof. Anthony D. Joseph

<http://cs162.eecs.Berkeley.edu>

## Review: Monitor with Condition Variables



- **Lock**: the lock provides mutual exclusion to shared data
  - Always acquire before accessing shared data structure
  - Always release after finishing with shared data
  - Lock initially free
- **Condition Variable**: a queue of threads waiting for something *inside* a critical section
  - Key idea: make it possible to go to sleep inside critical section by atomically releasing lock at time we go to sleep
  - Contrast to semaphores: Can't wait inside critical section

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.2

## Review: Condition Variables

- How do we change the `RemoveFromQueue()` routine to wait until something is on the queue?
  - Could do this by keeping a count of the number of things on the queue (with semaphores), but error prone
- **Condition Variable**: a queue of threads waiting for something *inside* a critical section
  - Key idea: allow sleeping inside critical section by atomically releasing lock at time we go to sleep
  - Contrast to semaphores: Can't wait inside critical section
- Operations:
  - **Wait(&lock)**: Atomically release lock and go to sleep. Re-acquire lock later, before returning.
  - **Signal()**: Wake up one waiter, if any
  - **Broadcast()**: Wake up all waiters
- Rule: Must hold lock when doing condition variable ops!
  - In Birrell paper, he says can perform `signal()` outside of lock – IGNORE HIM (this is only an optimization)

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.3

## Review: Mesa vs. Hoare Monitors

- Need to be careful about precise definition of signal and wait. Consider a piece of our dequeue code:
 

```
while (queue.isEmpty()) {
    dataready.wait(&lock); // If nothing, sleep
}
item = queue.dequeue(); // Get next item
```

  - Why didn't we do this?
 


```
if (queue.isEmpty()) {
    dataready.wait(&lock); // If nothing, sleep
}
item = queue.dequeue(); // Get next item
```
- Answer: depends on the type of scheduling
  - Hoare-style (most textbooks):
    - » Signaler gives lock, CPU to waiter; waiter runs immediately
    - » Waiter gives up lock, processor back to signaler when it exits critical section or if it waits again
  - Mesa-style (most real operating systems):
    - » Signaler keeps lock and processor
    - » Waiter placed on ready queue with no special priority
    - » **Practically, need to check condition again after wait**

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.4

## Extended Example: Readers/Writers Problem



- Motivation: Consider a shared database
  - Two classes of users:
    - » Readers – never modify database
    - » Writers – read and modify database
  - Is using a single lock on the whole database sufficient?
    - » Like to have many readers at the same time
    - » Only one writer at a time

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.5

## Basic Readers/Writers Solution

- Correctness Constraints:
  - Readers can access database when no writers
  - Writers can access database when no readers or writers
  - Only one thread manipulates state variables at a time
- Basic structure of a solution:
  - Reader()
    - Wait until no writers
    - Access data base
    - Check out – wake up a waiting writer
  - Writer()
    - Wait until no active readers or writers
    - Access database
    - Check out – wake up waiting readers or writer
  - State variables (Protected by a lock called “lock”):
    - » int AR: Number of active readers; initially = 0
    - » int WR: Number of waiting readers; initially = 0
    - » int AW: Number of active writers; initially = 0
    - » int WW: Number of waiting writers; initially = 0
    - » Condition okToRead = NIL
    - » Condition okToWrite = NIL

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.6

## Code for a Reader

```
Reader() {
    // First check self into system
    lock.Acquire();

    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;              // No. Writers exist
        okToRead.wait(&lock); // Sleep on cond var
        WR--;              // No longer waiting
    }

    AR++;                  // Now we are active!
    lock.release();

    // Perform actual read-only access
    AccessDatabase(ReadOnly);

    // Now, check out of system
    lock.Acquire();
    AR--;                  // No longer active
    if (AR == 0 && WW > 0) // No other active readers
        okToWrite.signal(); // Wake up one writer
    lock.Release();
}
```

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.7

## Code for a Writer

```
Writer() {
    // First check self into system
    lock.Acquire();

    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;              // No. Active users exist
        okToWrite.wait(&lock); // Sleep on cond var
        WW--;              // No longer waiting
    }

    AW++;                  // Now we are active!
    lock.release();

    // Perform actual read/write access
    AccessDatabase(ReadWrite);

    // Now, check out of system
    lock.Acquire();
    AW--;                  // No longer active
    if (WW > 0) {           // Give priority to writers
        okToWrite.signal(); // Wake up one writer
    } else if (WR > 0) {    // Otherwise, wake reader
        okToRead.broadcast(); // Wake all readers
    }
    lock.Release();
}
```

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.8

## Simulation of Readers/Writers solution

- Consider the following sequence of operators:
  - RI, R2, W1, R3
- On entry, each reader checks the following:

```
while ((AW + WW) > 0) { // Is it safe to read?
    WR++;                // No. Writers exist
    okToRead.wait(&lock); // Sleep on cond var
    WR--;                // No longer waiting
}
AR++;                    // Now we are active!
```
- First, R1 comes along:  
AR = 1, WR = 0, AW = 0, WW = 0
- Next, R2 comes along:  
AR = 2, WR = 0, AW = 0, WW = 0
- Now, readers make take a while to access database
  - Situation: Locks released
  - Only AR is non-zero

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.9

## Simulation(2)

- Next, W1 comes along:

```
while ((AW + AR) > 0) { // Is it safe to write?
    WW++;                // No. Active users exist
    okToWrite.wait(&lock); // Sleep on cond var
    WW--;                // No longer waiting
}
AW++;
```
- Can't start because of readers, so go to sleep:  
AR = 2, WR = 0, AW = 0, WW = 1
- Finally, R3 comes along:  
AR = 2, WR = 1, AW = 0, WW = 1
- Now, say that R2 finishes before R1:  
AR = 1, WR = 1, AW = 0, WW = 1
- Finally, last of first two readers (R1) finishes and wakes up writer:

```
if (AR == 0 && WW > 0) // No other active readers
    okToWrite.signal(); // Wake up one writer
```

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.10

## Simulation(3)

- When writer wakes up, get:  
AR = 0, WR = 1, AW = 1, WW = 0
- Then, when writer finishes:

```
if (WW > 0) { // Give priority to writers
    okToWrite.signal(); // Wake up one writer
} else if (WR > 0) { // Otherwise, wake reader
    okToRead.broadcast(); // Wake all readers
}
```

  - Writer wakes up reader, so get:  
AR = 1, WR = 0, AW = 0, WW = 0
- When reader completes, we are finished

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.11

## Questions

- Can readers starve? Consider Reader() entry code:

```
while ((AW + WW) > 0) { // Is it safe to read?
    WR++;                // No. Writers exist
    okToRead.wait(&lock); // Sleep on cond var
    WR--;                // No longer waiting
}
AR++;                    // Now we are active!
```
- What if we erase the condition check in Reader exit?

```
AR--;                    // No longer active
if (AR == 0 && WW > 0) // No other active readers
    okToWrite.signal(); // Wake up one writer
```
- Further, what if we turn the signal() into broadcast()

```
AR--;                    // No longer active
okToWrite.broadcast();   // Wake up one writer
```
- Finally, what if we use only one condition variable (call it "okToContinue") instead of two separate ones?
  - Both readers and writers sleep on this variable
  - Must use broadcast() instead of signal()

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.12

## Administrivia

---

- Midterm coming up soon
  - Wednesday 3/9 6-7:30PM
  - 10 EVANS (Seats: 237); 155 DWINELLE (Seats: 481)
    - » We will assign you to a room
  - Closed book, no calculators, one double-side page of handwritten notes
- No class that day, extra office hours
- Topics will include the material through lecture 12 (Wed 3/2)
  - Includes lectures, project 1, homeworks, readings, textbook

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.13

## BREAK

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.14

## Can we Construct Monitors from Semaphores?

---

- Locking aspect is easy: Just use a mutex
- Can we implement condition variables this way?

```
Wait() { semaphore.P(); }
Signal() { semaphore.V(); }
```

- Does this work better?

```
Wait(Lock lock) {
    lock.Release();
    semaphore.P();
    lock.Acquire();
}
Signal() { semaphore.V(); }
```

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.15

## Construction of Monitors from Semaphores (con't)

---

- Problem with previous try:
  - P and V are commutative – result is the same no matter what order they occur
  - Condition variables are NOT commutative
- Does this fix the problem?

```
Wait(Lock lock) {
    lock.Release();
    semaphore.P();
    lock.Acquire();
}
Signal() {
    if semaphore queue is not empty
        semaphore.V();
}
```

  - Not legal to look at contents of semaphore queue
  - There is a race condition – signaler can slip in after lock release and before waiter executes semaphore.P()
- It is actually possible to do this correctly
  - Complex solution for Hoare scheduling in book
  - Can you come up with simpler Mesa-scheduled solution?

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.16

## Monitor Conclusion

- Monitors represent the logic of the program
  - Wait if necessary
  - Signal when change something so any waiting threads can proceed
- Basic structure of monitor-based program:

```

lock
while (need to wait) {
    condvar.wait();
}
unlock

do something so no need to wait

lock

condvar.signal();

unlock
    
```

Check and/or update state variables  
Wait if necessary

Check and/or update state variables

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.17


## C-Language Support for Synchronization

- C language: Pretty straightforward synchronization
  - Just make sure you know *all* the code paths out of a critical section

```

int Rtn() {
    lock.acquire();
    ...
    if (exception) {
        lock.release();
        return errReturnCode;
    }
    ...
    lock.release();
    return OK;
}
    
```

- Watch out for `setjmp/longjmp`!
  - Can cause a non-local jump out of procedure
  - In example, procedure E calls `longjmp`, popping stack back to procedure B
  - If Procedure C had `lockacquire`, problem!



2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.18

## C++ Language Support for Synchronization

- Languages with exceptions like C++
  - Languages that support exceptions are problematic (easy to make a non-local exit without releasing lock)
  - Consider:

```

void Rtn() {
    lock.acquire();
    ...
    DoFoo();
    ...
    lock.release();
}
void DoFoo() {
    ...
    if (exception) throw errException;
    ...
}
    
```

- Notice that an exception in `DoFoo()` will exit without releasing the lock!

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.19

## C++ Language Support for Synchronization (con't)

- Must catch all exceptions in critical sections
  - Catch exceptions, release lock, and re-throw exception:

```

void Rtn() {
    lock.acquire();
    try {
        ...
        DoFoo();
        ...
    } catch (...) { // catch exception
        lock.release(); // release lock
        throw; // re-throw the exception
    }
    lock.release();
}
void DoFoo() {
    ...
    if (exception) throw errException;
    ...
}
    
```

- Even Better: `auto_ptr<T>` facility. See C++ Spec.
  - Can deallocate/free lock regardless of exit method

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.20

## Java Language Support for Synchronization

- Java has explicit support for threads and thread synchronization
- Bank Account example:

```
class Account {
    private int balance;
    // object constructor
    public Account (int initialBalance) {
        balance = initialBalance;
    }
    public synchronized int getBalance() {
        return balance;
    }
    public synchronized void deposit(int amount) {
        balance += amount;
    }
}
```
- Every object has an associated lock which gets automatically acquired and released on entry and exit from a *synchronized* method

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.21

## Java Language Support for Synchronization (con't)

- Java also has *synchronized* statements:

```
synchronized (object) {
    ...
}
```
- Since every Java object has an associated lock, this type of statement acquires and releases the object's lock on entry and exit of the body

- Works properly even with exceptions:

```
synchronized (object) {
    ...
    DoFoo();
    ...
}
void DoFoo() {
    throw errException;
}
```

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.22

## Java Language Support for Synchronization (con't 2)

- In addition to a lock, every object has a **single** condition variable associated with it
  - How to wait inside a synchronization method or block:

```
» void wait(long timeout); // Wait for timeout
» void wait(long timeout, int nanoseconds); //variant
» void wait();
```
  - How to signal in a synchronized method or block:

```
» void notify(); // wakes up oldest waiter
» void notifyAll(); // like broadcast, wakes everyone
```
  - Condition variables can wait for a bounded length of time. This is useful for handling exception cases:

```
t1 = time.now();
while (!ATMRequest()) {
    wait (CHECKPERIOD);
    t2 = time.now();
    if (t2 - t1 > LONG_TIME) checkMachine();
}
```
  - Not all Java VMs equivalent!
    - » Different scheduling policies, not necessarily preemptive!

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.23

## Recall: Better Implementation of Locks by Disabling Interrupts

- Key idea: maintain a lock variable and impose mutual exclusion only during operations on that variable

```
int mylock = FREE;
Acquire(&mylock) – Wait until lock is free, then grab
Release(&mylock) – Unlock, waking up anyone waiting

Acquire(int *lock) {
    disable interrupts;
    if (*lock == BUSY) {
        put thread on wait queue;
        Go to sleep();
        // Enable interrupts?
    } else {
        *lock = BUSY;
        enable interrupts;
    }
}

Release(int *lock) {
    disable interrupts;
    if (anyone on wait queue) {
        take thread off wait queue
        Place on ready queue;
    } else {
        *lock = FREE;
        enable interrupts;
    }
}
```

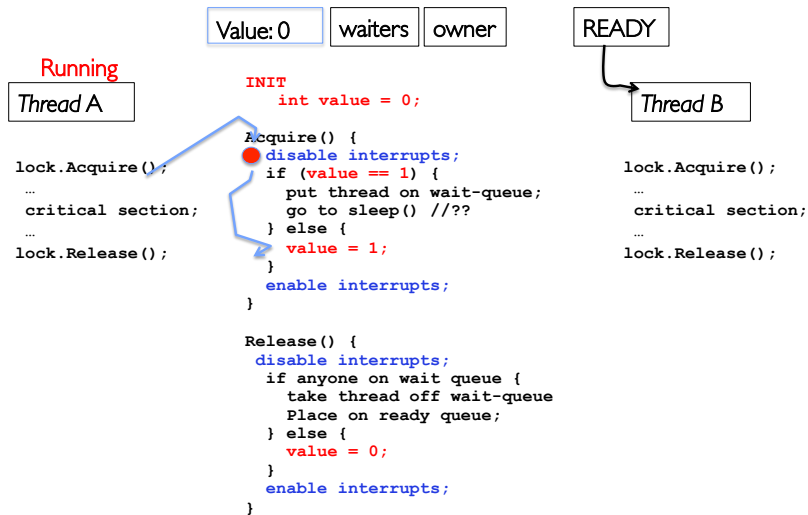
- Really only works in kernel – why?

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.24

## In-Kernel Lock Simulation

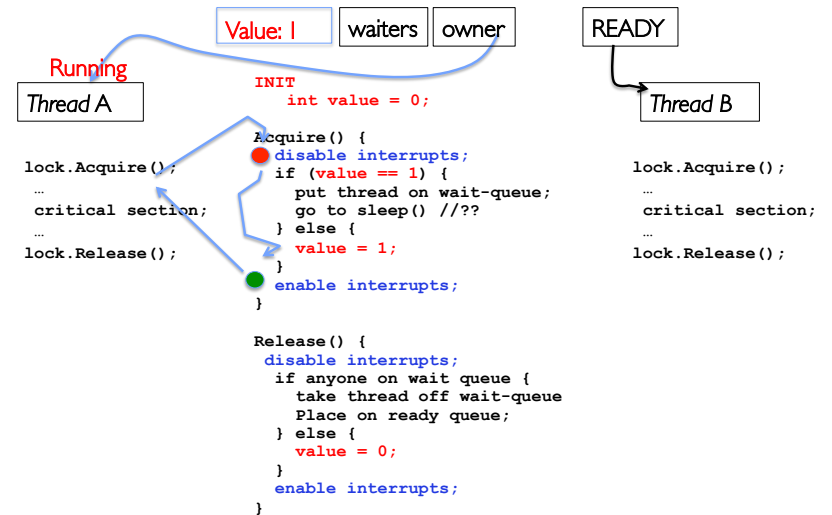


2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.25

## In-Kernel Lock Simulation

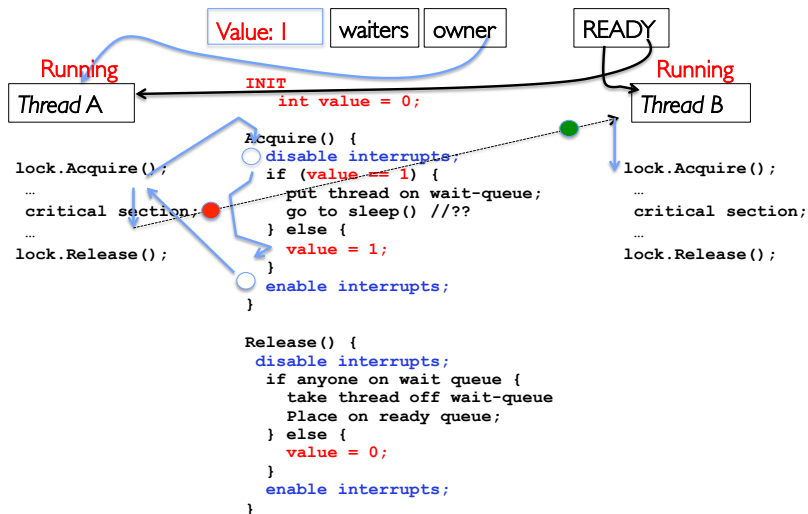


2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.26

## In-Kernel Lock Simulation

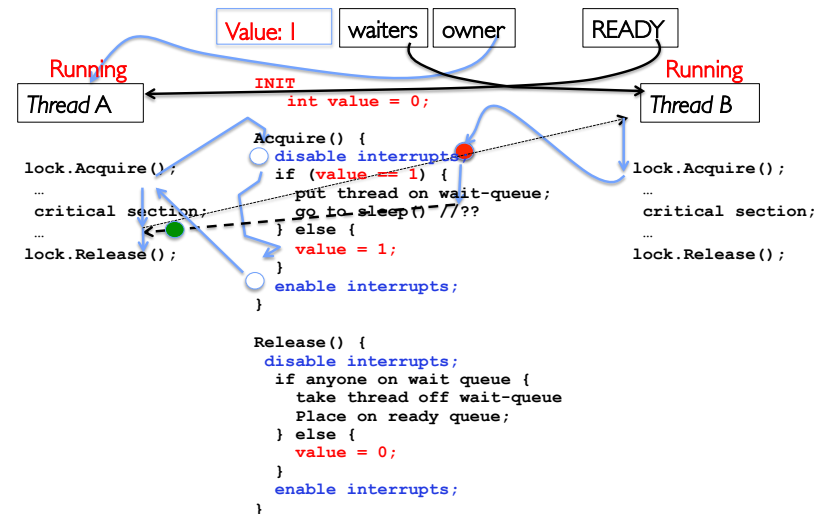


2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.27

## In-Kernel Lock Simulation

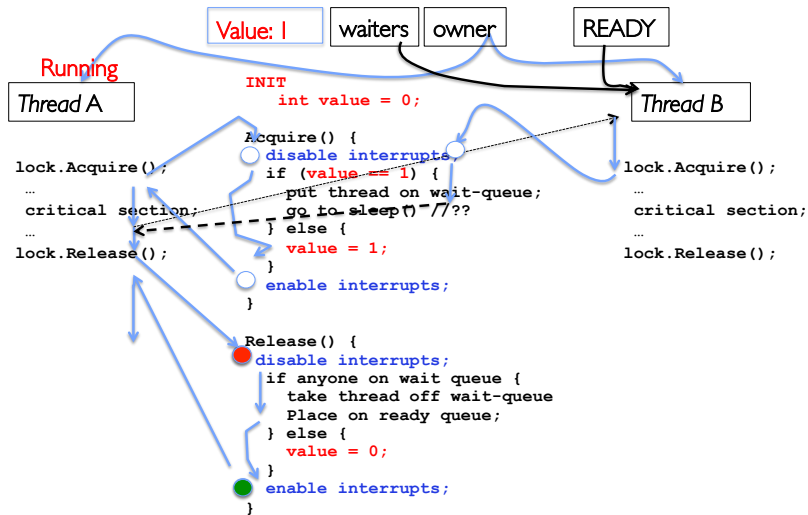


2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.28

## In-Kernel Lock: Simulation

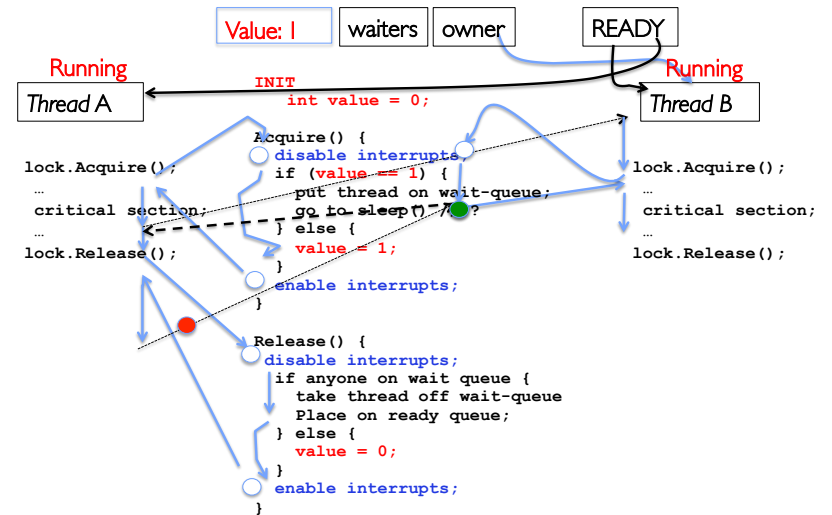


2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.29

## In-Kernel Lock: Simulation



2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.30

## Discussion

- Notice that Scheduling here involves deciding who to take off the wait queue
  - Could do by priority, etc.
- Same type of code works for in-kernel condition variables
  - The Wait queue becomes unique for each condition variable
  - Once again, transition to and from queues occurs with interrupts disabled

BREAK

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.31

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.32



## Synchronization Summary

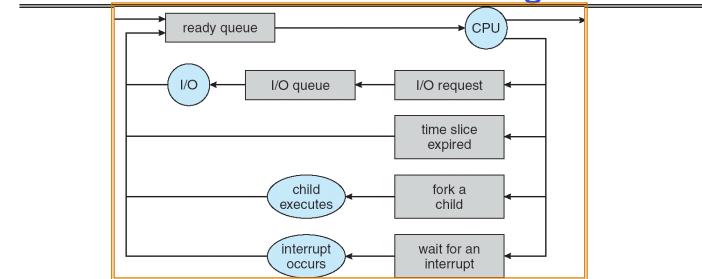
- **Semaphores**: Like integers with restricted interface
  - Two operations:
    - » **P()**: Wait if zero; decrement when becomes non-zero
    - » **V()**: Increment and wake a sleeping task (if exists)
    - » Can initialize value to any non-negative value
  - Use separate semaphore for each constraint
- **Monitors**: A lock plus zero or more condition variables
  - Always acquire lock before accessing shared data
  - Use condition variables to wait inside critical section
    - » Three Operations: **Wait()**, **Signal()**, **Broadcast()**

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.33

## Recall: CPU Scheduling



- Earlier, we talked about the life-cycle of a thread
  - Active threads work their way from Ready queue to Running to various waiting queues.
- Question: How is the OS to decide which of several tasks to take off a queue?
  - Obvious queue to worry about is ready queue
  - Others can be scheduled as well, however
- **Scheduling**: deciding which threads are given access to resources from moment to moment

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.34

## Recall: Scheduling Assumptions

- CPU scheduling big area of research in early 70's
- Many implicit assumptions for CPU scheduling:
  - One program per user
  - One thread per program
  - Programs are independent
- Clearly, these are unrealistic but they simplify the problem so it can be solved
  - For instance: is “fair” about fairness among users or programs?
    - » If I run one compilation job and you run five, you get five times as much CPU on many operating systems
- The high-level goal: Dole out CPU time to optimize some desired parameters of system

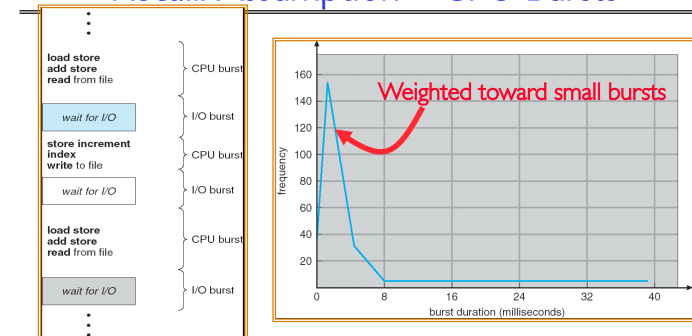


2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.35

## Recall: Assumption – CPU Bursts



- Execution model: programs alternate between bursts of CPU and I/O
  - Program typically uses the CPU for some period of time, then does I/O, then uses CPU again
  - Each scheduling decision is about which job to give to the CPU for use by its next CPU burst
  - With timeslicing, thread may be forced to give up CPU before finishing current CPU burst

2/22/16

Joseph CS162 ©UCB Spring 2016

Lec 9.36

## Scheduling Policy Goals/Criteria

- Minimize Response Time
  - Minimize elapsed time to do an operation (or job)
  - Response time is what the user sees:
    - » Time to echo a keystroke in editor
    - » Time to compile a program
    - » Real-time Tasks: Must meet deadlines imposed by World
- Maximize Throughput
  - Maximize operations (or jobs) per second
  - Throughput related to response time, but not identical:
    - » Minimizing response time will lead to more context switching than if you only maximized throughput
  - Two parts to maximizing throughput
    - » Minimize overhead (for example, context-switching)
    - » Efficient use of resources (CPU, disk, memory, etc)
- Fairness
  - Share CPU among users in some equitable way
  - Fairness is not minimizing average response time:
    - » Better average response time by making system less fair

2/22/16

Joseph CSI 62 ©UCB Spring 2016

Lec 9.37


## First-Come, First-Served (FCFS) Scheduling

- First-Come, First-Served (FCFS)
  - Also “First In, First Out” (FIFO) or “Run until done”
    - » In early systems, FCFS meant one program scheduled until done (including I/O)
    - » Now, means keep CPU until thread blocks
- Example:

Process	Burst Time
$P_1$	24
$P_2$	3
$P_3$	3



- Suppose processes arrive in the order:  $P_1, P_2, P_3$   
The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0; P_2 = 24; P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$
- Average Completion time:  $(24 + 27 + 30)/3 = 27$
- Convoy effect: short process behind long process

2/22/16

Joseph CSI 62 ©UCB Spring 2016

Lec 9.38

## FCFS Scheduling (Cont.)

- Example continued:
    - Suppose that processes arrive in order:  $P_2, P_3, P_1$   
Now, the Gantt chart for the schedule is:
- 
- | Process | Start Time | End Time | Completion Time |
|---------|------------|----------|-----------------|
| $P_2$   | 0          | 3        | 3               |
| $P_3$   | 3          | 6        | 6               |
| $P_1$   | 6          | 30       | 30              |
- Waiting time for  $P_1 = 6; P_2 = 0; P_3 = 3$
  - Average waiting time:  $(6 + 0 + 3)/3 = 3$
  - Average Completion time:  $(3 + 6 + 30)/3 = 13$
- In second case:
    - average waiting time is much better (before it was 17)
    - Average completion time is better (before it was 27)
  - FIFO Pros and Cons:
    - Simple (+)
    - Short jobs get stuck behind long ones (-)
      - » Safeway: Getting milk, always stuck behind cart full of small items. Upside: get to read about space aliens!

2/22/16

Joseph CSI 62 ©UCB Spring 2016

Lec 9.39

## Summary

- **Semaphores:** Like integers with restricted interface
  - Two operations:
    - » **P()**: Wait if zero; decrement when becomes non-zero
    - » **V()**: Increment and wake a sleeping task (if exists)
    - » Can initialize value to any non-negative value
  - Use separate semaphore for each constraint
- **Monitors:** A lock plus one or more condition variables
  - Always acquire lock before accessing shared data
  - Use condition variables to wait inside critical section
    - » Three Operations: **Wait()**, **Signal()**, and **Broadcast()**
- **Scheduling:** selecting a waiting process from the ready queue and allocating the CPU to it
- **FCFS Scheduling:**
  - Run threads to completion in order of submission
  - Pros: Simple
  - Cons: Short jobs get stuck behind long ones
- **Round-Robin Scheduling:**
  - Give each thread a small amount of CPU time when it executes; cycle between all ready threads
  - Pros: Better for short jobs
  - Cons: Poor when jobs are same length

2/22/16

Joseph CSI 62 ©UCB Spring 2016

Lec 9.40