


CS162  
Operating Systems and  
Systems Programming  
Lecture 14

Caching (Finished),  
Demand Paging

March 14<sup>th</sup>, 2016  
Prof. Anthony D. Joseph  
<http://cs162.eecs.Berkeley.edu>

Recall: In Machine Structures (eg. 61C) ...

- Caching is the key to memory system performance



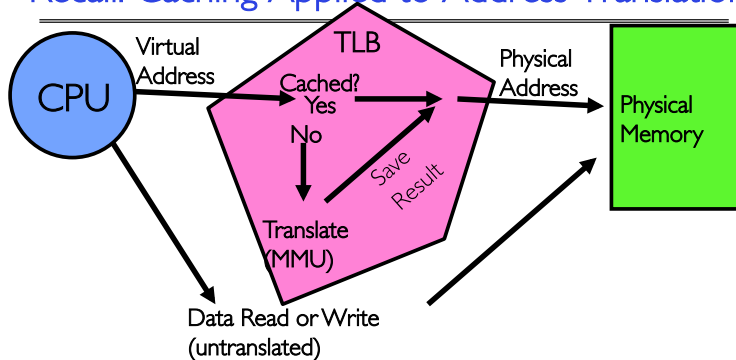
- Average Access time = (Hit Rate × HitTime) + (Miss Rate × MissTime)
- HitRate + MissRate = 1
- HitRate = 90% ⇒ Average Access Time = 19 ns
- HitRate = 99% ⇒ Average Access Time = 10.9ns

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.2

Recall: Caching Applied to Address Translation



- Question is one of page locality: does it exist?
  - Instruction accesses spend a lot of time on the same page (since accesses sequential)
  - Stack accesses have definite locality of reference
  - Data accesses have less page locality, but still some...
- Can we have a TLB hierarchy?
  - Sure: multiple levels at different sizes/speeds

Recall: What Actually Happens on a TLB Miss?

- Hardware traversed page tables:
  - On TLB miss, hardware in MMU looks at current page table to fill TLB (may walk multiple levels)
    - If PTE valid, hardware fills TLB and processor never knows
    - If PTE marked as invalid, causes Page Fault, after which kernel decides what to do afterwards
- Software traversed Page tables (like MIPS)
  - On TLB miss, processor receives TLB fault
  - Kernel traverses page table to find PTE
    - If PTE valid, fills TLB and returns from fault
    - If PTE marked as invalid, internally calls Page Fault handler
- Most chip sets provide hardware traversal
  - Modern operating systems tend to have more TLB faults since they use translation for many things
  - Examples:
    - shared segments
    - user-level portions of an operating system

3/14/16

Joseph CS162 ©UCB Spring 2016


Lec 14.3

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.4

## Transparent Exceptions: TLB/Page fault



- How to transparently restart faulting instructions?
  - (Consider load or store that gets TLB or Page fault)
  - Could we just skip faulting instruction?
    - » No: need to perform load or store after reconnecting physical page
- Hardware must help out by saving:
  - Faulting instruction and partial state
    - » Need to know which instruction caused fault
    - » Is single PC sufficient to identify faulting position????
  - Processor State: sufficient to restart user thread
    - » Save/restore registers, stack, etc
- What if an instruction has side-effects?

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.5

## Consider weird things that can happen

- What if an instruction has side effects?
  - Options:
    - » Unwind side-effects (easy to restart)
    - » Finish off side-effects (messy!)
  - Example 1: `mov (sp) +, 10`
    - » What if page fault occurs when write to stack pointer?
    - » Did `sp` get incremented before or after the page fault?
  - Example 2: `strcpy (r1), (r2)`
    - » Source and destination overlap: can't unwind in principle!
    - » IBM S/370 and VAX solution: execute twice – once read-only
- What about "RISC" processors?
  - For instance delayed branches?
    - » Example: `bne somewhere`  
`ld r1, (sp)`
    - » Precise exception state consists of two PCs: PC and nPC
  - Delayed exceptions:
    - » Example: `div r1, r2, r3`  
`ld r1, (sp)`
    - » What if takes many cycles to discover divide by zero, but load has already caused page fault?

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.6

## Precise Exceptions


- Precise  $\Rightarrow$  state of the machine is preserved as if program executed up to the offending instruction
  - All previous instructions **completed**
  - Offending instruction and all following instructions act **as if they have not even started**
  - Same system code will work on different implementations
  - Difficult in the presence of pipelining, out-of-order execution, ...
  - **MIPS takes this position**
- Imprecise  $\Rightarrow$  system software has to figure out what is where and put it all back together
- Performance goals often lead designers to forsake precise interrupts
  - system software developers, user, markets etc. usually wish they had not done this
- **Modern techniques for out-of-order execution and branch prediction help implement precise interrupts**

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.7

## Recall: TLB Organization



- Needs to be really fast
  - Critical path of memory access
    - » In simplest view: before the cache
    - » Thus, this adds to access time (reducing cache speed)
  - Seems to argue for Direct Mapped or Low Associativity
- However, needs to have very few conflicts!
  - With TLB, the Miss Time extremely high!
  - **This argues that cost of Conflict (Miss Time) is much higher than slightly increased cost of access (Hit Time)**
- **Thrashing:** continuous conflicts between accesses
  - What if use low order bits of page as index into TLB?
    - » First page of code, data, stack may map to same entry
    - » Need 3-way associativity at least?
  - What if use high order bits as index?
    - » TLB mostly unused for small programs

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.8

## Example: R3000 pipeline includes TLB “stages”

### MIPS R3000 Pipeline

Inst Fetch	Dcd/ Reg	ALU / E.A	Memory	Write Reg
TLB	I-Cache	RF	Operation	WB
		E.A.	TLB	D-Cache

### TLB

64 entry, on-chip, fully associative, software TLB fault handler

### Virtual Address Space

ASID	V. Page Number	Offset
6	20	12

0xx User segment (caching based on PT/TLB entry)  
 100 Kernel physical space, cached  
 101 Kernel physical space, uncached  
 11x Kernel virtual space

Allows context switching among  
 64 user processes without TLB flush

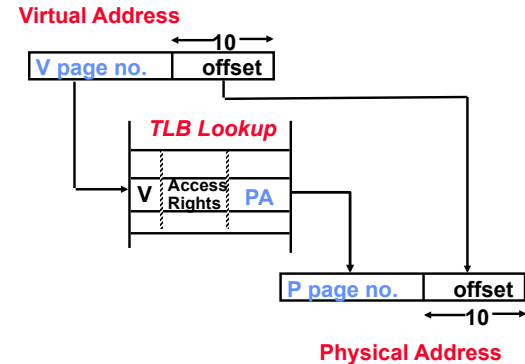
3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.9

## Reducing translation time further

- As described, TLB lookup is in serial with cache lookup:



- Machines with TLBs go one step further: they overlap TLB lookup with cache access.
  - Works because offset available early

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.10

## Overlapping TLB & Cache Access (1/2)

- Main idea:
  - Offset in virtual address exactly covers the “cache index” and “byte select”
  - Thus can select the cached byte(s) in parallel to perform address translation

virtual address 

Virtual Page #	Offset
----------------	--------

physical address 

tag / page #	index	byte
--------------	-------	------

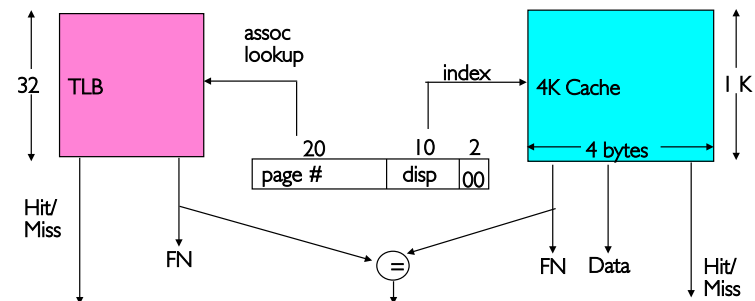
3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.11

## Overlapping TLB & Cache Access

- Here is how this might work with a 4K cache:



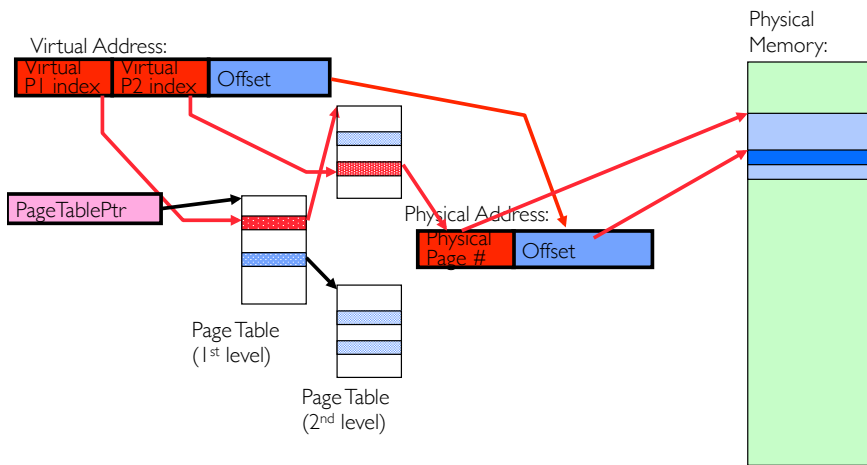
- What if cache size is increased to 8KB?
  - Overlap not complete
  - Need to do something else. See CS152/252
- Another option: **Virtual Caches**
  - Tags in cache are virtual addresses
  - Translation only happens on cache misses

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.12

## Putting Everything Together: Address Translation

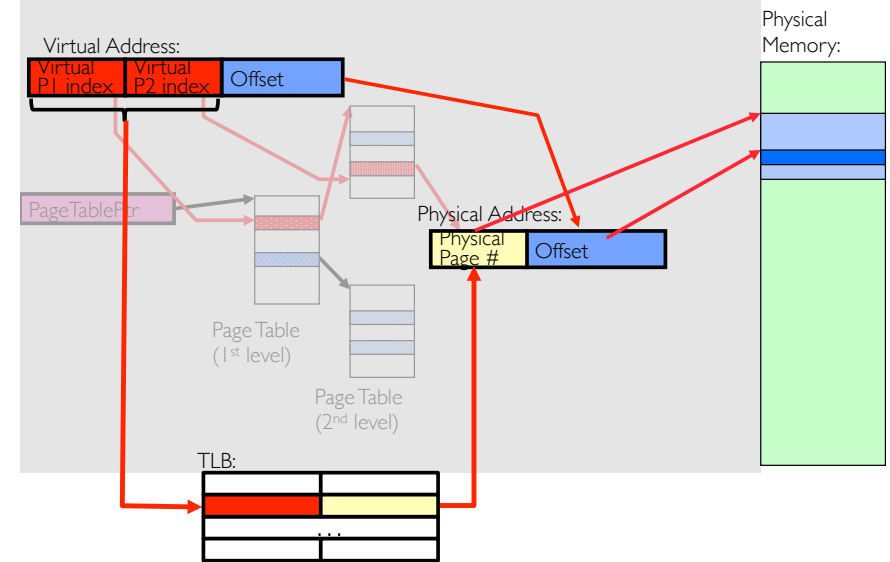


3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.13

## Putting Everything Together: TLB

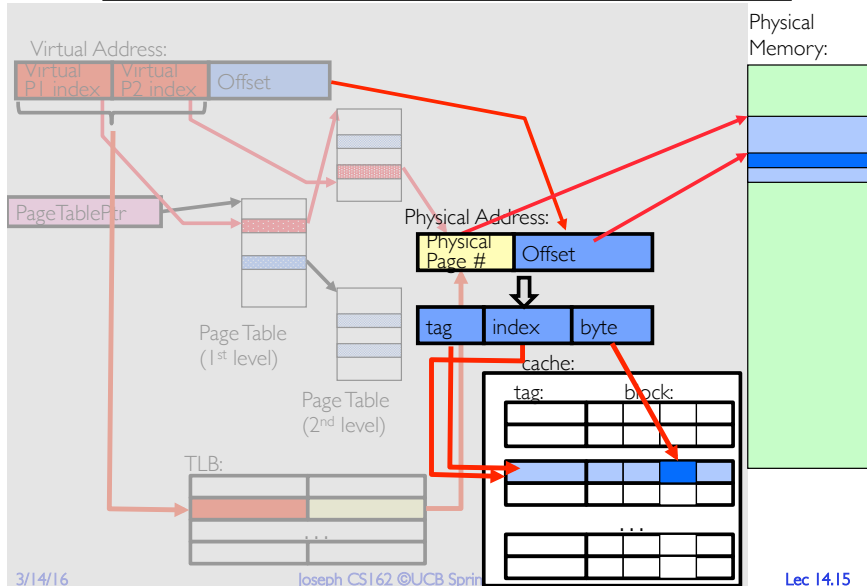


3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.14

## Putting Everything Together: Cache

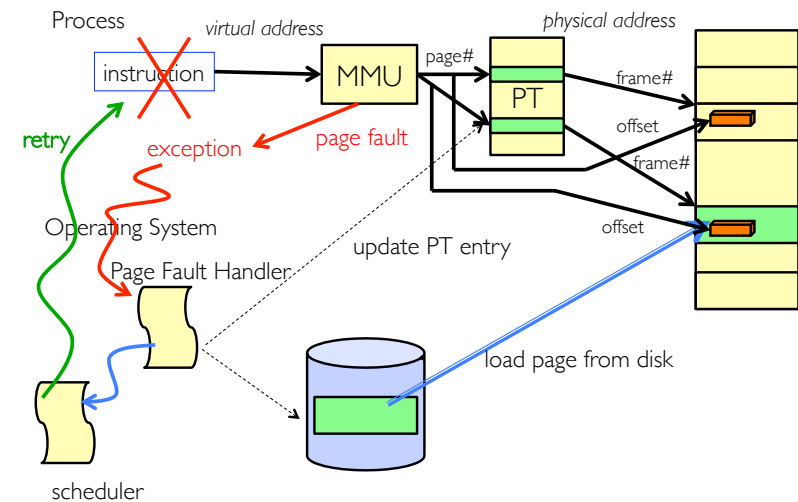


3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.15

## Next Up: What happens when ...




3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.16

## Administrivia

- Midterm #1 grades posted to gradescope and autograder



- Midterm #1 is 11% of total course grade
  - If your grade is >2 std devs below mean, talk with prof or staff
- Project 2 has been released
  - Get started early as design doc is due Monday 3/28

BREAK

## Where are all places that caching arises in OSes?


- Direct use of caching techniques
  - paged virtual memory (mem as cache for disk)
  - TLB (cache of PTEs)
  - file systems (cache disk blocks in memory)
  - DNS (cache hostname => IP address translations)
  - Web proxies (cache recently accessed pages)
- Which pages to keep in memory?
  - All-important “Policy” aspect of virtual memory
  - Will spend a bit more time on this in a moment

## Impact of caches on Operating Systems

- Indirect - dealing with cache effects (e.g., sync state across levels)
- Process scheduling
  - which and how many processes are active ?
  - large memory footprints versus small ones ?
  - priorities ?
  - Shared pages mapped into VAS of multiple processes ?
- Impact of thread scheduling on cache performance
  - rapid interleaving of threads (small quantum) may degrade cache performance
    - » increase average memory access time (AMAT) !!!
- Designing operating system data structures for cache performance
- Maintaining the correctness of various caches
  - TLB consistency:
    - » With PT across context switches ?
    - » Across updates to the PT ?

## Working Set Model

- As a program executes it transitions through a sequence of “working sets” consisting of varying sized subsets of the address space




3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.21

## Cache Behavior under WS model




- Amortized by fraction of time the Working Set is active
- Transitions from one WS to the next
- Capacity, Conflict, Compulsory misses
- Applicable to memory caches and pages. Others ?

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.22

## Another model of Locality: Zipf



- Likelihood of accessing item of rank  $r$  is  $\propto 1/r^a$
- Although rare to access items below the top few, there are so many that it yields a “heavy tailed” distribution
- Substantial value from even a tiny cache
- Substantial misses from even a very large one


3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.23

## Demand Paging


- Modern programs require a lot of physical memory
  - Memory per system growing faster than 25%-30%/year
- But they don't use all their memory all of the time
  - 90-10 rule: programs spend 90% of their time in 10% of their code
  - Wasteful to require all of user's code to be in memory
- Solution: use main memory as cache for disk



3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.24



- Disk is larger than physical memory  $\Rightarrow$ 
  - In-use virtual memory can be bigger than physical memory
  - Combined memory of running processes much larger than physical memory
    - » More programs fit into memory, allowing more concurrency
- Principle: **Transparent Level of Indirection** (page table)
  - Supports flexible placement of physical data
    - » Data could be on disk or somewhere across network
  - Variable location of data transparent to user program
    - » Performance issue, not correctness issue

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.25

## Demand Paging is Caching

- Since Demand Paging is Caching, must ask:
  - What is block size?
    - » 1 page
  - What is organization of this cache (i.e. direct-mapped, set-associative, fully-associative)?
    - » Fully associative: arbitrary virtual  $\rightarrow$  physical mapping
  - How do we find a page in the cache when look for it?
    - » First check TLB, then page-table traversal
  - What is page replacement policy? (i.e. LRU, Random...)
    - » This requires more explanation... (kinda LRU)
  - What happens on a miss?
    - » Go to lower level to fill miss (i.e. disk)
  - What happens on a write? (write-through, write back)
    - » Definitely write-back – need dirty bit!

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.26

## Recall: What is in a Page Table Entry

- What is in a Page Table Entry (or PTE)?
  - Pointer to next-level page table or to actual page
  - Permission bits: valid, read-only, read-write, write-only
- Example: Intel x86 architecture PTE:
  - Address same format previous slide (10, 10, 12-bit offset)
  - Intermediate page tables called “Directories”

Page Frame Number (Physical Page Number)	Free (OS)	0	L	D	A	PCD	PWT	U	W	P
31-12	11-9	8	7	6	5	4	3	2	1	0

**P:** Present (same as “valid” bit in other architectures)

**W:** Writeable

**U:** User accessible

**PWT:** Page write transparent: external cache write-through

**PCD:** Page cache disabled (page cannot be cached)

**A:** Accessed: page has been accessed recently

**D:** Dirty (PTE only): page has been modified recently

**L:** L=1  $\Rightarrow$  4MB page (directory only).

Bottom 22 bits of virtual address serve as offset

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.27

## Demand Paging Mechanisms

- PTE helps us implement demand paging
  - Valid  $\Rightarrow$  Page in memory, PTE points at physical page
  - Not Valid  $\Rightarrow$  Page not in memory; use info in PTE to find it on disk when necessary
- Suppose user references page with invalid PTE?
  - Memory Management Unit (MMU) traps to OS
    - » Resulting trap is a “Page Fault”
  - What does OS do on a Page Fault?:
    - » Choose an old page to replace
    - » If old page modified (“D=1”), write contents back to disk
    - » Change its PTE and any cached TLB to be invalid
    - » Load new page into memory from disk
    - » Update page table entry, invalidate TLB for new entry
    - » Continue thread from original faulting location
  - TLB for new page will be loaded when thread continued!
  - While pulling pages off disk for one process, OS runs another process from ready queue
    - » Suspended process sits on wait queue



3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.28


BREAK

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.29

## Loading an executable into memory



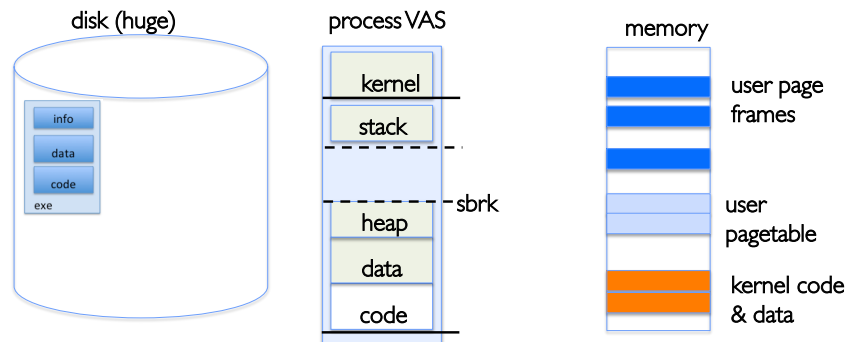
- .exe
  - lives on disk in the file system
  - contains contents of code & data segments, relocation entries and symbols
  - OS loads it into memory, initializes registers (and initial stack pointer)
  - program sets up stack and heap upon initialization: crt0

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.30

## Create Virtual Address Space of the Process



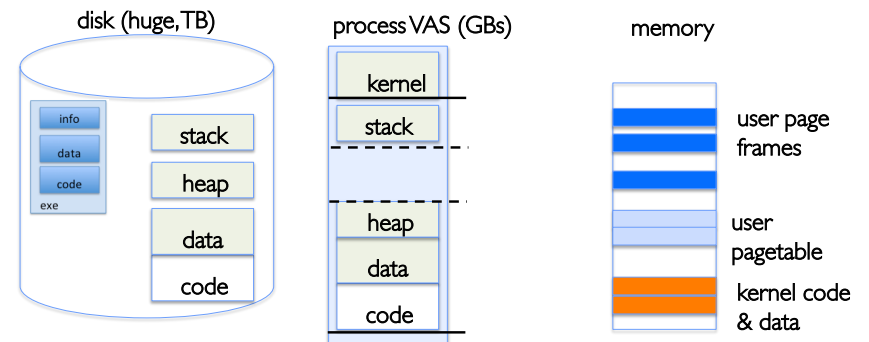
- Utilized pages in the VAS are backed by a page block on disk
  - Called the backing store or swap file
  - Typically in an optimized block store, but can think of it like a file

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.31

## Create Virtual Address Space of the Process



- User Page table maps entire VAS
- All the utilized regions are backed on disk
  - swapped into and out of memory as needed
- For every process

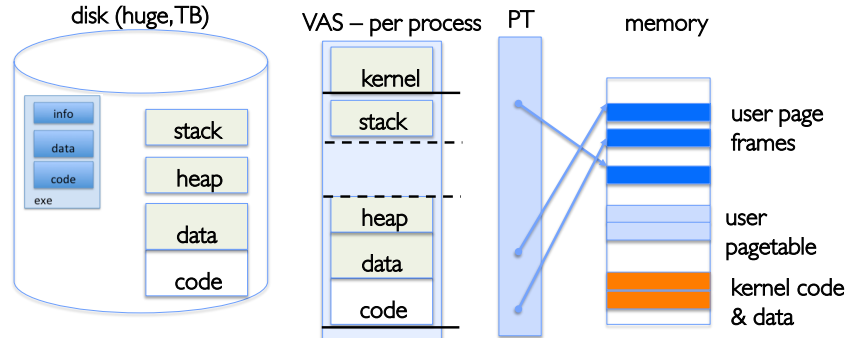
3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.32



## Create Virtual Address Space of the Process




- User Page table maps entire VAS
  - Resident pages to the frame in memory they occupy
  - The portion of it that the HW needs to access must be resident in memory

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.33

## Provide Backing Store for VAS



- User Page table maps entire VAS
- Resident pages mapped to memory frames
- For all other pages, OS must record where to find them on disk

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.34

## What data structure maps non-resident pages to disk?

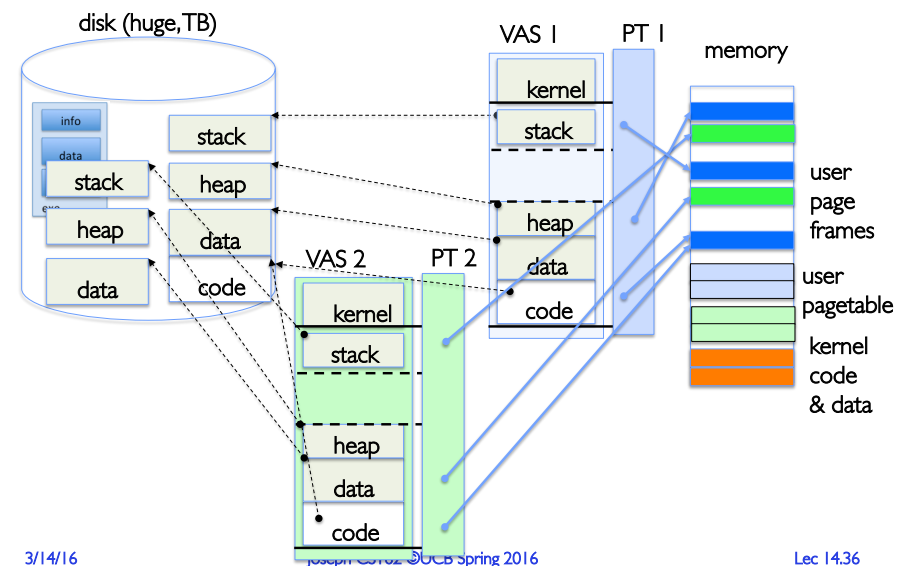
- FindBlock(PID, page#) => disk\_block
  - Some OSs utilize spare space in PTE for paged blocks
  - Like the PT, but purely software
- Where to store it?
  - In memory – can be compact representation if swap storage is contiguous on disk
  - Could use hash table (like Inverted PT)
- Usually want backing store for resident pages too
- May map code segment directly to on-disk image
  - Saves a copy of code to swap file
- May share code segment with multiple instances of the program

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.35

## Provide Backing Store for VAS

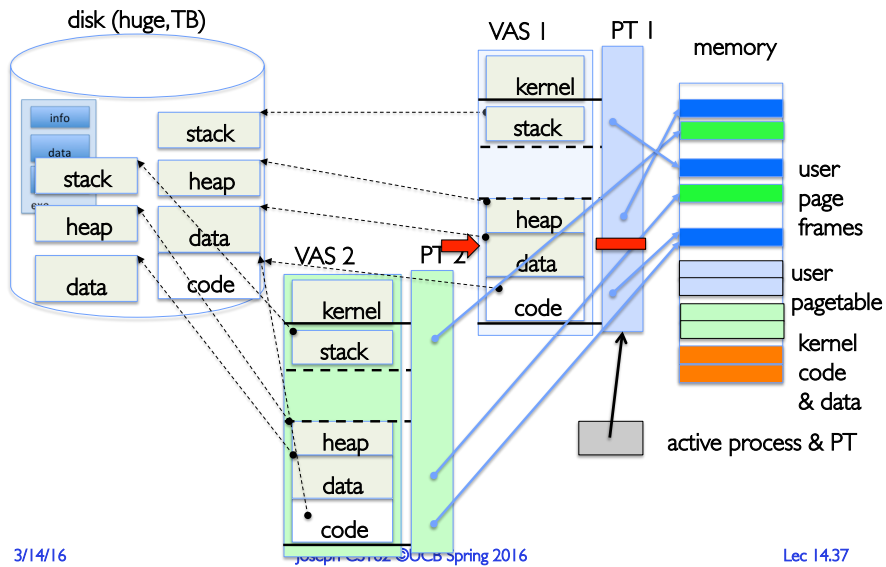


3/14/16

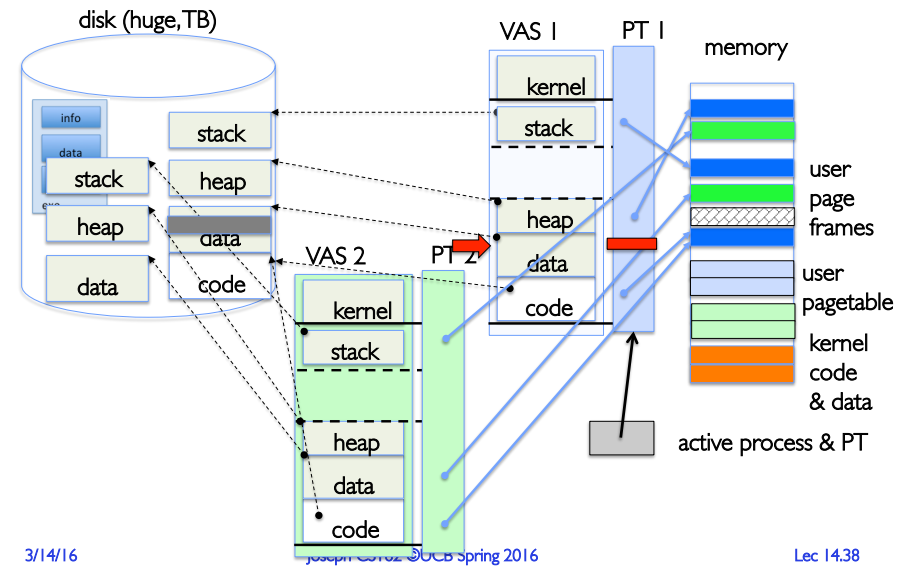
Joseph CS162 ©UCB Spring 2016

Lec 14.36

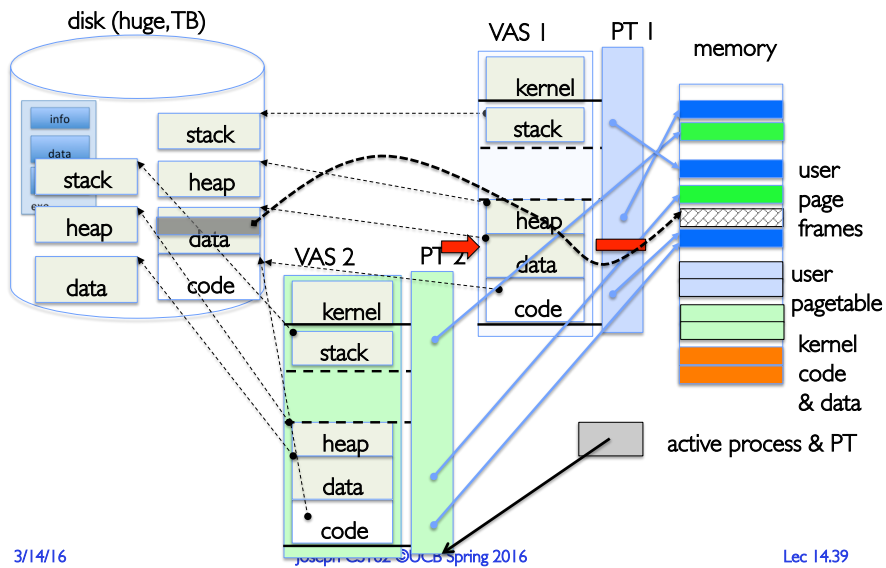
## On page Fault ...



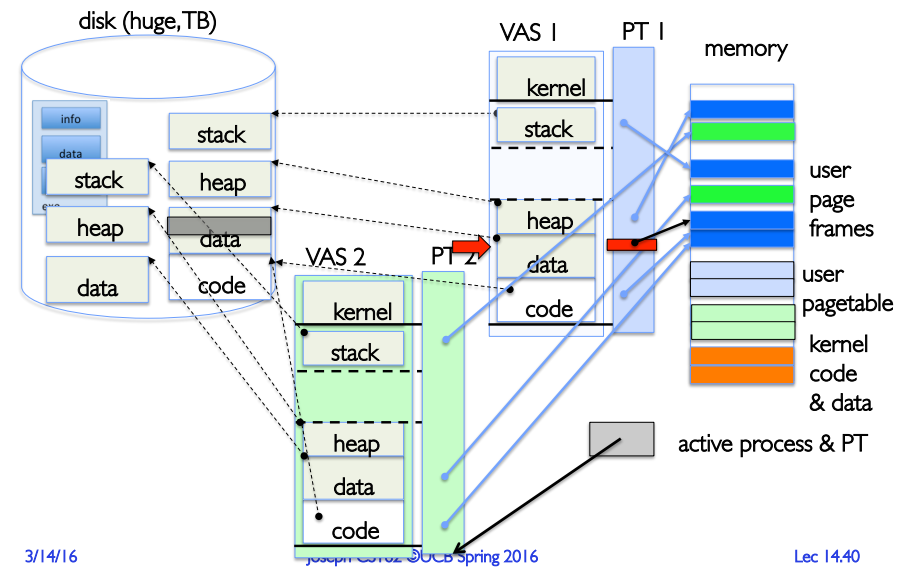
## On page Fault ... find & start load



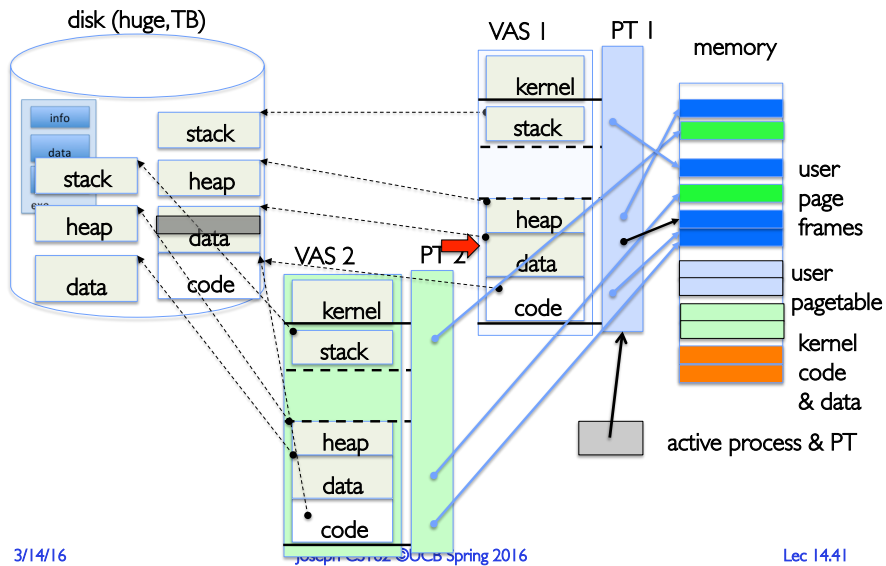
## On page Fault ... schedule other P or T



## On page Fault ... update PTE



## Eventually reschedule faulting thread

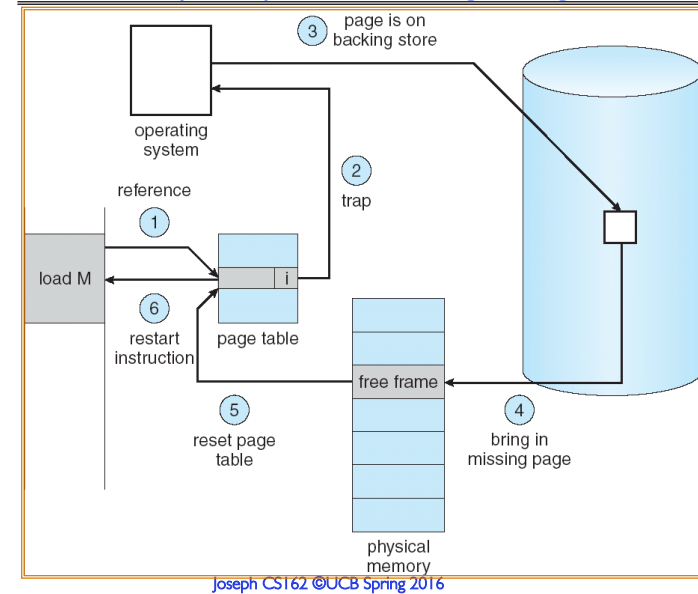


3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.41

## Summary: Steps in Handling a Page Fault

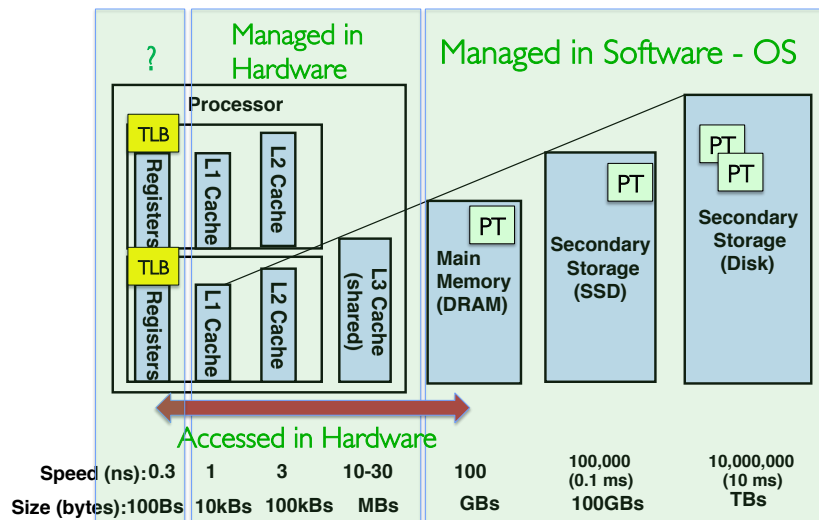


3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.42

## Management & Access to the Memory Hierarchy



3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.43

## Some following questions

- During a page fault, where does the OS get a free frame?
  - Keeps a free list
  - Unix runs a “reaper” if memory gets too full
  - As a last resort, evict a dirty page first
- How can we organize these mechanisms?
  - Work on the replacement policy
- How many page frames/process?
  - Like thread scheduling, need to “schedule” memory resources:
    - » utilization? fairness? priority?
  - allocation of disk paging bandwidth

3/14/16

Joseph CS162 ©UCB Spring 2016

Lec 14.44

## Summary

---

- A cache of translations called a “Translation Lookaside Buffer” (TLB)
  - Relatively small number of entries ( $< 512$ )
  - Fully Associative (Since conflict misses expensive)
  - TLB entries contain PTE and optional process ID
- On TLB miss, page table must be traversed
  - If located PTE is invalid, cause Page Fault
- On context switch/change in page table
  - TLB entries must be invalidated somehow
- TLB is logically in front of cache
  - Thus, needs to be overlapped with cache access to be really fast
- Precise Exception specifies a single instruction for which:
  - All previous instructions have completed (committed state)
  - No following instructions nor actual instruction have started