

CS162 Operating Systems and Systems Programming Lecture 18

File Systems

April 4th, 2016
Prof. Anthony D. Joseph
<http://cs162.eecs.Berkeley.edu>

Quick Aside: Big Projects

- What is a big project?
 - Time/work estimation is hard
 - Programmers are eternal optimists (it will only take two days!)
 - » This is why we bug you about starting the project early
 - » Had a grad student who used to say he just needed "10 minutes" to fix something. Two hours later...
- Can a project be efficiently partitioned?
 - Partitionable task decreases in time as you add people
 - But, if you require communication:
 - » Time reaches a minimum bound
 - » With complex interactions, time increases!
 - Mythical person-month problem:
 - » You estimate how long a project will take
 - » Starts to fall behind, so you add more people
 - » Project takes even more time!



4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.2

Techniques for Partitioning Tasks

- Functional
 - Person A implements threads, Person B implements semaphores, Person C implements locks...
 - Problem: Lots of communication across APIs
 - » If B changes the API, A may need to make changes
 - » Story: Large airline company spent \$200 million on a new scheduling and booking system. Two teams "working together." After two years, went to merge software. Failed! Interfaces had changed (documented, but no one noticed). Result: would cost another \$200 million to fix.
- Task
 - Person A designs, Person B writes code, Person C tests
 - May be difficult to find right balance, but can focus on each person's strengths (Theory vs systems hacker)
 - Since Debugging is hard, Microsoft has two testers for each programmer
- Most CS162 project teams are functional, but people have had success with task-based divisions


4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.3

Communication

- More people mean more communication
 - Changes have to be propagated to more people
 - Think about person writing code for most fundamental component of system: everyone depends on them!
 - *You should be meeting in person at least twice/week!*
- Miscommunication is common
 - "Index starts at 0? I thought you said 1!"
- Who makes decisions?
 - Individual decisions are fast but trouble
 - Group decisions take time
 - Centralized decisions require a big picture view (someone who can be the "system architect")
- Often designating someone as system architect can be a good thing
 - Better not be clueless
 - Better have good people skills
 - Better let other people do work



4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.4

Coordination



- More people \Rightarrow no one can make all meetings!
 - They miss decisions and associated discussion
 - Example from earlier class: one person missed meetings and did something group had rejected
- People have different work styles
 - Some people work in the morning, some at night
 - How do you decide when to meet or work together?
- What about project slippage?
 - It will happen, guaranteed!
 - Example: phase 4 of one project, everyone busy but not talking. One person way behind. No one knew until very end – too late!
- Hard to add people to existing group
 - Members have already figured out how to work together

Recall: Device Drivers

- **Device Driver:** Device-specific code in the kernel that interacts directly with the device hardware
 - Supports a standard, internal interface
 - Same kernel I/O system can interact easily with different device drivers
 - Special device-specific configuration supported with the `ioctl()` system call
- Device Drivers typically divided into two pieces:
 - Top half: accessed in call path from system calls
 - » implements a set of **standard, cross-device calls** like `open()`, `close()`, `read()`, `write()`, `ioctl()`, `strategy()`
 - » This is the kernel's interface to the device driver
 - » Top half will start I/O to device, may put thread to sleep until finished
 - Bottom half: run as interrupt routine
 - » Gets input or transfers next block of output
 - » May wake sleeping threads if I/O now complete

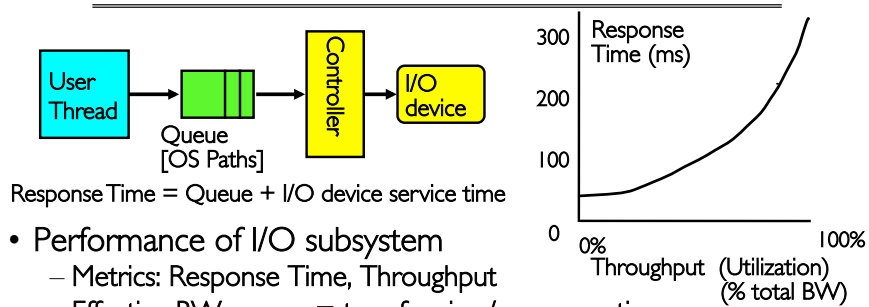
Kernel vs User-level I/O

- Both are popular and practical for different reasons:
- **Kernel-level drivers** for critical devices that must keep running (e.g., display drivers)
 - Programming is a major effort, correct operation of the rest of the kernel depends on correct driver operation
- **User-level drivers** for devices that are non-threatening (e.g., USB devices in Linux: `libusb`)
 - Provide higher-level primitives to the programmer, avoid every driver doing low-level I/O register tweaking
 - The multitude of USB devices can be supported by Less-Than-Wizard programmers
 - New drivers don't have to be compiled for each version of the OS, and loaded into the kernel

Kernel vs User-level Programming Styles

- **Kernel-level drivers**
 - Have a much more limited set of resources available:
 - » Only a fraction of `libc` routines typically available
 - » Memory allocation (e.g., Linux `kmalloc`) much more limited in capacity and required to be physically contiguous
 - » Should avoid blocking calls
 - » Can use asynchrony with other kernel functions but tricky with user code
- **User-level drivers**
 - Similar to other application programs but:
 - » Will be called often – should do its work fast, or postpone it – or do it in the background
 - » Can use threads, blocking operations (usually much simpler) or non-blocking or asynchronous

Recall: I/O Performance



- Performance of I/O subsystem

- Metrics: Response Time, Throughput
- Effective BW per op = transfer size / response time
 - » $\text{EffectiveBW}(n) = n / (S + n/B) = B / (1 + SB/n)$
- Contributing factors to latency:
 - » Software paths (can be loosely modeled by a queue)
 - » Hardware controller
 - » I/O device service time

- Queuing behavior:


- Can lead to big increases in latency as utilization increases

4/4/16 – Solutions?

Joseph CS162 ©UCB Spring 2016

Lec 18.9

Performance: Multiple Outstanding Requests



- Suppose each read takes 10 ms to service
- If a process works for 100 ms after each read, what is the utilization of the disk?
 - $U = 10 \text{ ms} / 110 \text{ ms} = 9\%$
- What if there are two such processes?
 - $U = (10 \text{ ms} + 10 \text{ ms}) / 120 \text{ ms} \approx 17\%$
- What if each of those processes have two such threads?

4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.10

Recall: How do we Hide I/O Latency?

- **Blocking Interface:** “Wait”

- When request data (e.g., `read()` system call), put process to sleep until data is ready
- When write data (e.g., `write()` system call), put process to sleep until device is ready for data

- **Non-blocking Interface:** “Don’t Wait”

- Returns quickly from read or write request with count of bytes successfully transferred to kernel
- Read may return nothing, write may write nothing

- **Asynchronous Interface:** “Tell Me Later”

- When requesting data, take pointer to user’s buffer, return immediately; later kernel fills buffer and notifies user
- When sending data, take pointer to user’s buffer, return immediately; later kernel takes data and notifies user

4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.11

Administrivia

- HW4 – Releases on Monday 4/11 (Due 4/25)
- Project 2 code due next Friday
- Midterm II: Coming up in 2.5 weeks! (4/20)
 - 6-7:30PM (aa-eh 10 Evans, ej-oe 155 Dwinelle)
 - Covers lectures #13 to 21
 - 1 page of hand-written notes, both sides

4/4/16

Joseph CS162 ©UCB Spring 2016


Lec 18.12

BREAK

Recall: Building a File System

- **File System:** Layer of OS that transforms block interface of disks (or other block devices) into Files, Directories, etc.
- File System Components
 - Disk Management: collecting disk blocks into files
 - Naming: Interface to find files by name, not by blocks
 - Protection: Layers to keep data secure
 - Reliability/Durability: Keeping of files durable despite crashes, media failures, attacks, etc.
- User vs. System View of a File
 - User's view:
 - » Durable Data Structures
 - System's view (system call interface):
 - » Collection of Bytes (UNIX)
 - » Doesn't matter to system what kind of data structures you want to store on disk!
 - System's view (inside OS):
 - » Collection of blocks (a block is a logical transfer unit, while a sector is the physical transfer unit)
 - » Block size \geq sector size; in UNIX, block size is 4KB

Recall: Translating from User to System View



- What happens if user says: give me bytes 2—12?
 - Fetch block corresponding to those bytes
 - Return just the correct portion of the block
- What about: write bytes 2—12?
 - Fetch block
 - Modify portion
 - Write out Block
- Everything inside File System is in whole size blocks
 - For example, `getc()`, `putc()` \Rightarrow buffers something like 4096 bytes, even if interface is one byte at a time
- From now on, file is a collection of blocks

So You Are Going to Design a File System ...

- What factors are critical to the design choices?
- Durable data store \Rightarrow it's all on disk
- (Hard) Disks Performance !!!
 - Maximize sequential access, minimize seeks
- Open before Read/Write
 - Can perform protection checks and look up where the actual file resource are, in advance
- Size is determined as they are used !!!
 - Can write (or read zeros) to expand the file
 - Start small and grow, need to make room
- Organized into directories
 - What data structure (on disk) for that?
- Need to allocate / free blocks
 - Such that access remains efficient

Recall: Disk Management Policies

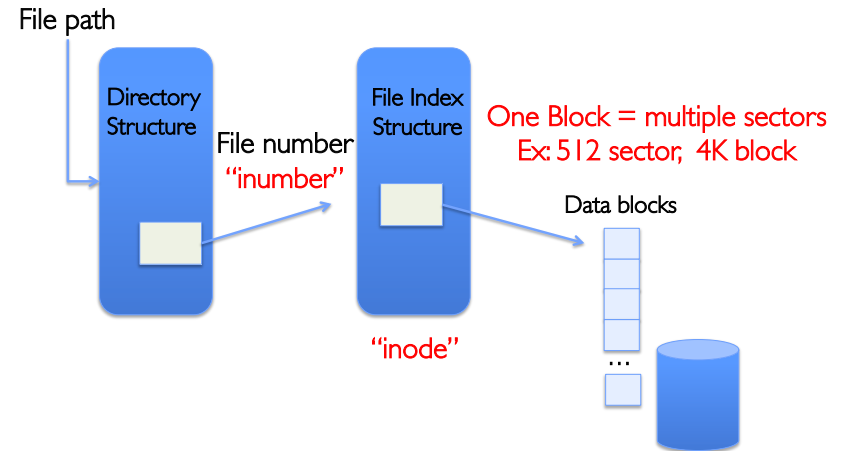
- Basic entities on a disk:
 - File**: user-visible group of blocks arranged sequentially in logical space
 - Directory**: user-visible index mapping names to files (next lecture)
- Access disk as linear array of sectors. Two Options:
 - Identify sectors as vectors [cylinder, surface, sector], sort in cylinder-major order, not used anymore
 - Logical Block Addressing (LBA)**: Every sector has integer address from zero up to max number of sectors
 - Controller translates from address \Rightarrow physical position
 - First case: OS/BIOS must deal with bad sectors
 - Second case: hardware shields OS from structure of disk
- Need way to track free disk blocks
 - Link free blocks together \Rightarrow too slow today
 - Use bitmap to represent free space on disk
- Need way to structure files: **File Header**
 - Track which blocks belong at which offsets within the logical file structure
 - Optimize placement of files' disk blocks to match access and usage patterns**

4/4/16

Joseph CSI 62 ©UCB Spring 2016

Lec 18.17

Components of a File System



4/4/16

Joseph CSI 62 ©UCB Spring 2016

Lec 18.18

Components of a file system

file name $\xrightarrow{\text{offset}}$ file number $\xrightarrow{\text{offset}}$ Storage block
 directory index structure

- Open performs **Name Resolution**
 - Translates pathname into a “file number”
 - Used as an “index” to locate the blocks
 - Creates a file descriptor in PCB within kernel
 - Returns a “handle” (another integer) to user process
- Read, Write, Seek, and Sync operate on handle
 - Mapped to descriptor and to blocks

4/4/16

Joseph CSI 62 ©UCB Spring 2016

Lec 18.19

Directories

FAVORITES	Name	Date Modified	Size	Kind
culer	bse	Yesterday, 6:21 PM	--	Folder
All My Files	Classes	Oct 13, 2014, 10:19 PM	--	Folder
AirDrop	AIIT2008	Oct 13, 2014, 10:11 PM	--	Folder
Applications	CS-Scholars	Oct 13, 2014, 10:11 PM	--	Folder
Desktop	cs61ci-f08	Oct 13, 2014, 10:17 PM	--	Folder
Documents	cs61ci-f09	Oct 13, 2014, 10:19 PM	--	Folder
Downloads	cs162	Today, 8:36 AM	--	Folder
	AndersonDahlin	Oct 13, 2014, 10:11 PM	--	Folder
	fa14	Today, 8:36 AM	--	Folder
	162prereqcheckSept8.xlsx	Sep 10, 2014, 3:20 PM	36 KB	Micros...kbook
	coursecomparison.xlsx	Aug 6, 2014, 7:50 AM	31 KB	Micros...kbook
	CS 162 apps.xlsx	Jun 29, 2014, 6:35 AM	53 KB	Micros...kbook
	cs162git	Sep 23, 2014, 11:33 AM	--	Folder
	devel	Oct 15, 2014, 11:40 AM	--	Folder
	exams	Oct 13, 2014, 10:12 PM	--	Folder
	gitprojects	Oct 8, 2014, 4:52 PM	--	Folder
	group0	Today, 8:35 AM	--	Folder
	pintos	Today, 8:35 AM	--	Folder
	src	Today, 8:35 AM	--	Folder
	gradesheet.xls	Sep 19, 2014, 4:48 PM	68 KB	Micros...kbook
	CSI Section Coverage.xlsx	Aug 22, 2014, 1:29 PM	11 KB	Micros...kbook
	Lectures	Today, 8:22 AM	--	Folder
	pintos-notes.txt	Sep 14, 2014, 2:10 PM	1 KB	Plain Text
	pintos.pdf	Jul 21, 2014, 10:17 AM	549 KB	PDF Document
	roster-9-13.xls	Sep 13, 2014, 5:12 PM	83 KB	Micros...kbook
	roster-9-19.xls	Sep 19, 2014, 4:39 PM	84 KB	Micros...kbook
	staff.xlsx	Aug 6, 2014, 7:14 AM	34 KB	Micros...kbook
	student	Oct 13, 2014, 10:12 PM	--	Folder
	studentsExcelFile-10-20	Yesterday, 9:53 AM	84 KB	Micros...kbook
	syllabus-fa14.xlsx	Sep 12, 2014, 10:00 AM	38 KB	Micros...kbook
	tmp	Oct 13, 2014, 10:12 PM	--	Folder
	pintos	Aug 8, 2014, 6:06 AM	--	Folder
	sp14	May 14, 2014, 9:02 PM	--	Folder
	cs194	Oct 13, 2014, 10:16 PM	--	Folder
	cs262b	Aug 7, 2013, 7:55 AM	--	Folder

4/4/16

Joseph CSI 62 ©UCB Spring 2016

Lec 18.20

Directory

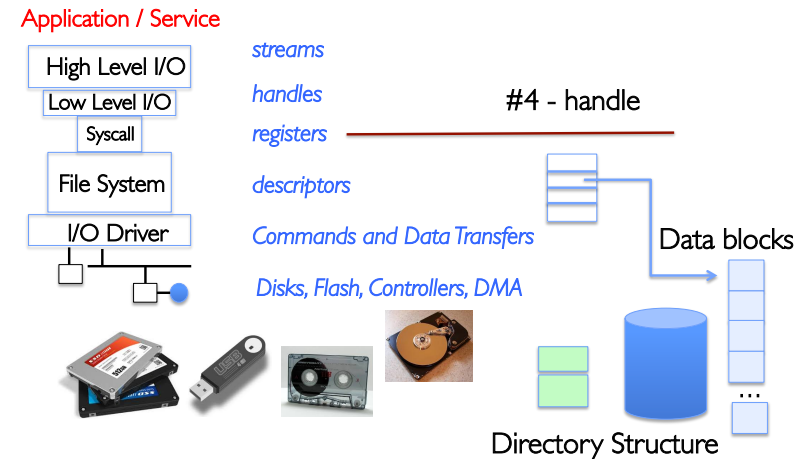
- Basically a hierarchical structure
- Each directory entry is a collection of
 - Files
 - Directories
 - » A link to another entries
- Each has a name and attributes
 - Files have data
- Links (hard links) make it a DAG, not just a tree
 - Softlinks (aliases) are another name for an entry

4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.21

I/O & Storage Layers



4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.22

File


- Named permanent storage
 - Contains
 - Data
 - » Blocks on disk somewhere
 - Metadata (Attributes)
 - » Owner, size, last opened, ...
 - » Access rights
 - R, W, X
 - Owner, Group, Other (in Unix systems)
 - Access control list in Windows system
- Diagram illustrating the components of a file:
- Data blocks:** Represented by a vertical stack of blue blocks.
 - File handle:** A horizontal line representing the handle to the file.
 - File descriptor:** A label for the handle.
 - Fileobject (inode):** A label for the handle.
 - Position:** A label for the handle.

4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.23

In-Memory File System Structures




- Open system call:
 - Resolves file name, finds file control block (inode)
 - Makes entries in per-process and system-wide tables
 - Returns index (called "file handle") in open-file table

4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.24

In-Memory File System Structures



- Read/write system calls:
 - Use file handle to locate inode
 - Perform appropriate reads or writes

4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.25

BREAK

4/4/16

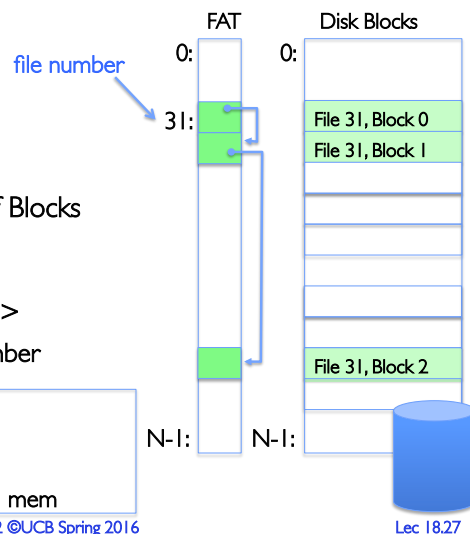
Joseph CS162 ©UCB Spring 2016

Lec 18.26

Our first filesystem: FAT (File Allocation Table)

- The most commonly used filesystem in the world!

- Assume (for now) we have a way to translate a path to a "file number"
 - i.e., a directory structure
- Disk Storage is a collection of Blocks
 - Just hold file data
- Example: file_read 31, < 2, x >
 - Index into FAT with file number
 - Follow linked list to block
 - Read the block from disk into memory



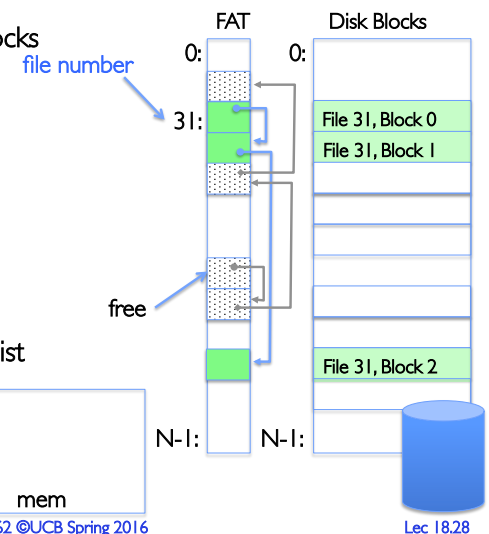
4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.27

FAT Properties

- File is collection of disk blocks
- FAT is linked list 1-1 with blocks
- File Number is index of root of block list for the file
- File offset ($o = B \cdot x$)
- Follow list to get block #
- Unused blocks \Leftrightarrow FAT free list



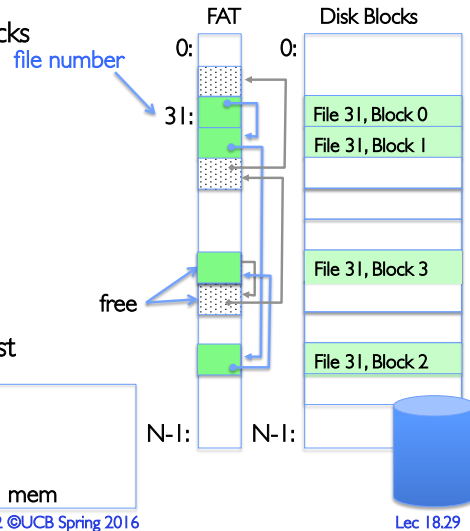
4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.28

FAT Properties

- File is collection of disk blocks
- FAT is linked list 1-I with blocks
- File Number is index of root of block list for the file
- File offset ($o = B \times x$)
- Follow list to get block #
- Unused blocks \leftrightarrow FAT free list
- Ex: `file_write(51, <3, y>)`
 - Grab blocks from free list
 - Linking them into file



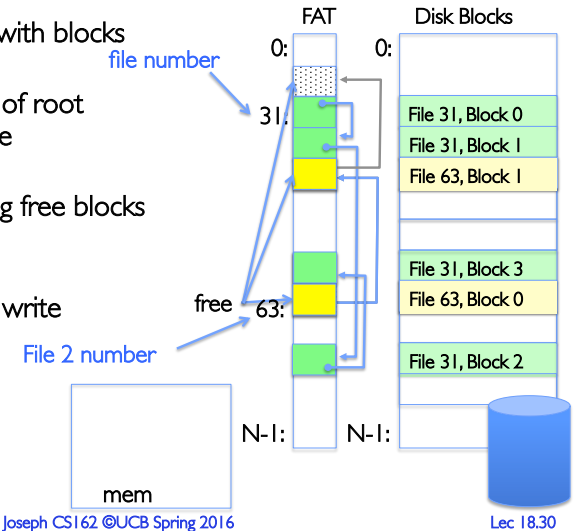
4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.29

FAT Properties

- File is collection of disk blocks
- FAT is linked list 1-I with blocks
- File Number is index of root of block list for the file
- Grow file by allocating free blocks and linking them in
- Ex: Create file, write, write



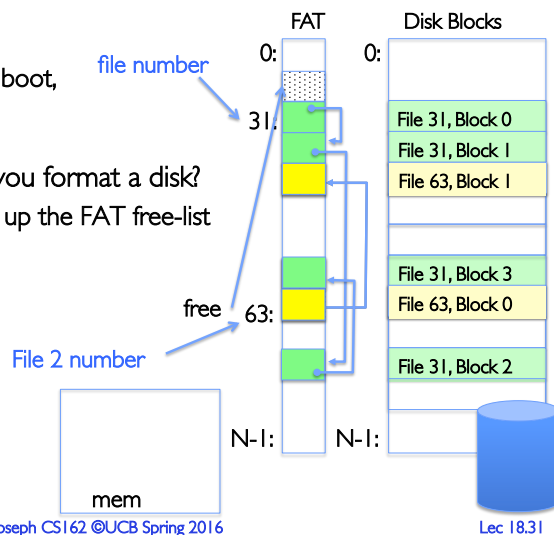
4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.30

FAT Assessment

- *FAT32 (32 instead of 12 bits) used in Windows, USB drives, SD cards, ...*
- Where is FAT stored?
 - On Disk, restore on boot, copy in memory
- What happens when you format a disk?
 - Zero the blocks, link up the FAT free-list
- Simple




4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.31

FAT Assessment

- Time to find block (large files) ??
- Block layout for file ???
- Sequential Access ???
- Random Access ???
- Fragmentation ???
- Small files ???
- Big files ???




4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.32

What about the Directory?



- Essentially a file containing
<file_name: file_number> mappings
- Free space for new entries
- In FAT: attributes kept in directory (!!!)
- Each directory a linked list of entries
- Where do you find root directory ("/")?

4/4/16

Joseph CSI 62 ©UCB Spring 2016

Lec 18.33

Directory Structure (cont'd)

- How many disk accesses to resolve "/my/book/count"?
 - Read in file header for root (fixed spot on disk)
 - Read in first data block for root
 - » Table of file name/index pairs. Search linearly – ok since directories typically very small
 - Read in file header for "my"
 - Read in first data block for "my"; search for "book"
 - Read in file header for "book"
 - Read in first data block for "book"; search for "count"
 - Read in file header for "count"
- **Current working directory:** Per-address-space pointer to a directory (inode) used for resolving file names
 - Allows user to specify relative filename instead of absolute path (say CWD="/my/book" can resolve "count")

4/4/16

Joseph CSI 62 ©UCB Spring 2016

Lec 18.34

Many Huge FAT Security Holes!

- FAT has no access rights
- FAT has no header in the file blocks
- Just gives an index into the FAT
 - (file number = block number)

4/4/16

Joseph CSI 62 ©UCB Spring 2016

Lec 18.35

Characteristics of Files

- Most files are small
- Most of the space is occupied by the rare big ones

A Five-Year Study of File-System Metadata

NITIN AGRAWAL
University of Wisconsin, Madison
and
WILLIAM J. BOLOSKEY, JOHN R. DOUCEUR, and JACOB R. LORCH
Microsoft Research




Fig. 2. Histograms of files by size.

Fig. 4. Histograms of bytes by containing file size.


4/4/16

Joseph CSI 62 ©UCB Spring 2016

Lec 18.36

So What About a “Real” File System?

• Meet the inode:



4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.37

Summary

- File System:
 - Transforms blocks into Files and Directories
 - Optimize for access and usage patterns
 - Maximize sequential access, allow efficient random access
- File (and directory) defined by header, called “inode”
- File Allocation Table (FAT) Scheme
 - Linked-list approach
 - Very widely used: Cameras, USB drives, SD cards
 - Simple to implement, but poor performance

4/4/16

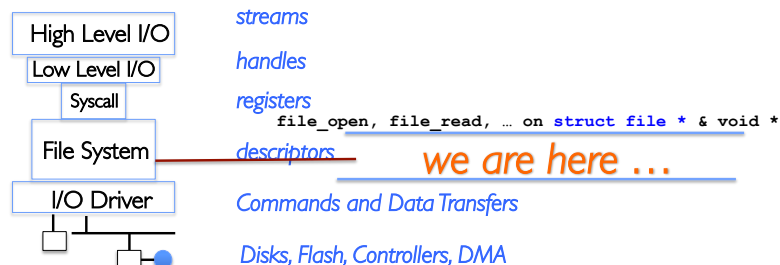
Joseph CS162 ©UCB Spring 2016

Lec 18.38

I/O & Storage Layers

Operations, Entities and Interface

Application / Service



4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.39

Recall: C Low level I/O

- Operations on File Descriptors – as OS object representing the state of a file
 - User has a “handle” on the descriptor

```
#include <fcntl.h>
#include <unistd.h>
#include <sys/types.h>

int open (const char *filename, int flags [, mode_t mode])
int creat (const char *filename, mode_t mode)
int close (int filedes)
```

Bit vector of:

- Access modes (Rd, Wr, ...)
- Open Flags (Create, ...)
- Operating modes (Appends, ...)

Bit vector of Permission Bits:

- User|Group|Other X R|W|X

http://www.gnu.org/software/libc/manual/html_node/Opening-and-Closing-Files.html

4/4/16

Joseph CS162 ©UCB Spring 2016

Lec 18.40

Recall: C Low Level Operations

```
ssize_t read (int fildes, void *buffer, size_t maxsize)
- returns bytes read, 0 => EOF, -1 => error
ssize_t write (int fildes, const void *buffer, size_t size)
- returns bytes written
```

```
off_t lseek (int fildes, off_t offset, int whence)
```

```
int fsync (int fildes) – wait for i/o to finish
```

```
void sync (void) – wait for ALL to finish
```

- When write returns, data is on its way to disk and can be read, but it may not actually be permanent!