

CSI 62  
Operating Systems and  
Systems Programming  
Lecture 24

Security

April 27<sup>th</sup>, 2016  
Prof. Anthony D. Joseph  
<http://cs162.eecs.Berkeley.edu>

## What is Computer Security Today?

- Computing in the presence of an adversary!
  - Adversary is the security field's defining characteristic
- Reliability, robustness, and fault tolerance
  - Dealing with Mother Nature (random failures)
- Security
  - Dealing with actions of a knowledgeable attacker dedicated to causing harm
  - Surviving malice, and not just mischance
- Wherever there is an adversary, there is a computer security problem!



4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.2

## Protection vs. Security

- **Protection**: mechanisms for controlling access of programs, processes, or users to resources
  - Page table mechanism
  - Round-robin schedule
  - Data encryption
- **Security**: use of protection mech. to prevent misuse of resources
  - Misuse defined with respect to policy
    - » E.g.: prevent exposure of certain sensitive information
    - » E.g.: prevent unauthorized modification/deletion of data
  - Need to consider external environment the system operates in
    - » Most well-constructed system cannot protect information if user accidentally reveals password – social engineering challenge

4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.3

## Security Requirements

- Authentication
  - Ensures that a user is who is claiming to be
- Data integrity
  - Ensure that data is not changed from source to destination or after being written on a storage device
- Confidentiality
  - Ensures that data is read only by authorized users
- Non-repudiation
  - Sender/client can't later claim didn't send/write data
  - Receiver/server can't claim didn't receive/write data

4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.4

## Securing Communication: Cryptography

- Cryptography: communication in the presence of adversaries
- Studied for thousands of years
  - See the Simon Singh's **The Code Book** for an excellent, highly readable history
- Central goal: confidentiality
  - How to encode information so that an adversary can't extract it, but a friend can
- General premise: there is a key, possession of which allows decoding, but without which decoding is infeasible
  - Thus, key must be kept secret and not guessable


4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.5

## Using Symmetric Keys

- Same key for encryption and decryption
- Achieves confidentiality
- Vulnerable to tampering and replay attacks




4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.6

## Symmetric Keys

- Can just XOR plaintext with the key
  - Easy to implement, but easy to break using frequency analysis
  - Unbreakable alternative: XOR with one-time pad
    - » Use a different key for each message



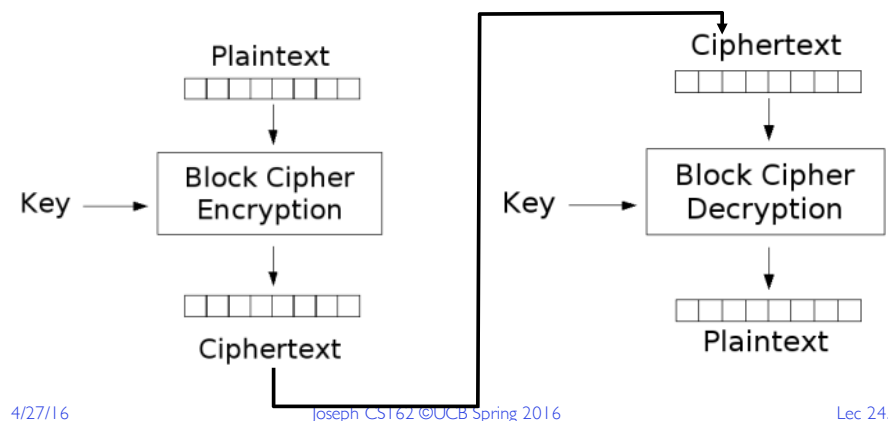
4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.7

## Block Ciphers with Symmetric Keys

- More sophisticated (e.g., block cipher) algorithms
  - Works with a block size (e.g., 64 bits)
- Can encrypt blocks separately:
  - Same plaintext same ciphertext
- Much better:
  - Add in counter and/or link ciphertext of previous block



4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.8

## Symmetric Key Ciphers - DES & AES

- Data Encryption Standard (DES)
  - Developed by IBM in 1970s, standardized by NBS/NIST
  - 56-bit key (decreased from 64 bits at NSA's request)
  - Still fairly strong other than brute-forcing the key space
    - » But custom hardware can crack a key in < 24 hours
  - Today many financial institutions use Triple DES
    - » DES applied 3 times, with 3 keys totaling 168 bits
- Advanced Encryption Standard (AES)
  - Replacement for DES standardized in 2002
  - Key size: 128, 192 or 256 bits
- How fundamentally strong are they?
  - No one knows (no proofs exist)


4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.9

## Authentication in Distributed Systems

- What if identity must be established across network?



- Need way to prevent exposure of information while still proving identity to remote system
- Many of the original UNIX tools sent passwords over the wire “in clear text”
  - » E.g.: telnet, ftp, yp (yellow pages, for distributed login)
  - » Result: Snooping programs widespread
- What do we need? Cannot rely on physical security!
  - Encryption: Privacy, restrict receivers
  - Authentication: Remote Authenticity, restrict senders


4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.10

## Authentication via Secret Key

- Main idea: entity proves identity by decrypting a secret encrypted with its own key
  - K – secret key shared only by A and B
- A can ask B to authenticate itself by decrypting a nonce, i.e., random value, x
  - Avoid replay attacks (attacker impersonating client or server)
- Vulnerable to man-in-the-middle attack




4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.11

## Secure Hash Function



- Hash Function: Short summary of data (message)
  - For instance,  $h_1 = H(M_1)$  is the hash of message  $M_1$ 
    - »  $h_1$  fixed length, despite size of message  $M_1$
    - » Often,  $h_1$  is called the “digest” of  $M_1$
- Hash function  $H$  is considered secure if
  - It is infeasible to find  $M_2$  with  $h_1 = H(M_2)$ ; i.e., can't easily find other message with same digest as given message
  - It is infeasible to locate two messages,  $m_1$  and  $m_2$ , which “collide”, i.e. for which  $H(m_1) = H(m_2)$
  - A small change in a message changes many bits of digest/can't tell anything about message given its hash

4/27/16


Joseph CSI 62 ©UCB Spring 2016

Lec 24.12

## Integrity: Cryptographic Hashes

- Basic building block for integrity: cryptographic hashing
  - Associate hash with byte-stream, receiver verifies match
    - » Assures data hasn't been modified, either accidentally – or maliciously
- Approach:
  - Sender computes a secure digest of message  $m$  using  $H(x)$ 
    - »  $H(x)$  is a publicly known hash function
    - » Digest  $d = \text{HMAC}(K, m) = H(K \parallel H(K \parallel m))$
    - »  $\text{HMAC}(K, m)$  is a hash-based message authentication function
  - Send digest  $d$  and message  $m$  to receiver
  - Upon receiving  $m$  and  $d$ , receiver uses shared secret key,  $K$ , to recompute  $\text{HMAC}(K, m)$  and see whether result agrees with  $d$

## Using Hashing for Integrity



Can encrypt  $m$  for confidentiality

## Standard Cryptographic Hash Functions

- MD5 (Message Digest version 5)
  - Developed in 1991 (Rivest), produces 128 bit hashes
  - Widely used (RFC 1321)
  - Broken (1996-2008): attacks that find collisions
- SHA-1 (Secure Hash Algorithm)
  - Developed in 1995 (NSA) as MD5 successor with 160 bit hashes
  - Widely used (SSL/TLS, SSH, PGP, IPSEC)
  - Broken in 2005, government use discontinued in 2010
- SHA-2 (2001)
  - Family of SHA-224, SHA-256, SHA-384, SHA-512 functions
- HMAC's are secure even with older "insecure" hash functions

## Administrivia


- Final Exam
  - Monday May 9th 3-6 PM Wheeler Auditorium
  - All material from the course
    - » With slightly more focus on second half, but you are still responsible for all the material
  - Two double-sided sheets of handwritten notes
  - No calculators
- Review session
  - May 6<sup>th</sup> 1-4 PM 390 Hearst Mining

BREAK

## Key Distribution

- How do you get shared secret to both places?
  - For instance: how do you send authenticated, secret mail to someone who you have never met?
  - Must negotiate key over private channel
    - » Exchange code book
    - » Key cards/memory stick/others
- Third Party: Authentication Server (like Kerberos)
  - Notation:
    - »  $K_{xy}$  is key for talking between  $x$  and  $y$
    - »  $(\dots)^K$  means encrypt message  $(\dots)$  with the key  $K$
    - » Clients:  $A$  and  $B$ , Authentication server  $S$
  - $A$  asks server for key:
    - »  $A \rightarrow S$ : [Hi! I'd like a key for talking between  $A$  and  $B$ ]
    - » Not encrypted. Others can find out if  $A$  and  $B$  are talking
  - Server returns session key encrypted using  $B$ 's key
    - »  $S \rightarrow A$ : **Message** [ Use  $K_{ab}$  (This is  $A$ ! Use  $K_{ab}$ ) <sup>$K_{sb}$</sup>  ] <sup>$K_{sa}$</sup>
    - » This allows  $A$  to know, " $S$  said use this key"
  - Whenever  $A$  wants to talk with  $B$ 
    - »  $A \rightarrow B$ : **Ticket** [ This is  $A$ ! Use  $K_{ab}$  ] <sup>$K_{sb}$</sup>
    - » Now,  $B$  knows that  $K_{ab}$  is sanctioned by  $S$

## Authentication Server Continued [Kerberos]




- Details
  - Both  $A$  and  $B$  use passwords (shared with key server) to decrypt return from key servers
  - Add in timestamps to limit how long tickets will be used to prevent attacker from replaying messages later
  - Also have to include encrypted checksums (hashed version of message) to prevent malicious user from inserting things into messages/changing messages
  - Want to minimize # times  $A$  types in password
    - »  $A \rightarrow S$  (Give me temporary secret)
    - »  $S \rightarrow A$  (Use  $K_{temp-sa}$  for next 8 hours) <sup>$K_{sa}$</sup>
    - » Can now use  $K_{temp-sa}$  in place of  $K_{sa}$  in protocol

## Asymmetric Encryption (Public Key)

- Idea: use two different keys, one to encrypt ( $e$ ) and one to decrypt ( $d$ )
  - A **key pair**
- Crucial property: knowing  $e$  does not give away  $d$
- Therefore  $e$  can be public: everyone knows it!
- If Alice wants to send to Bob, she fetches Bob's public key (say from Bob's home page) and encrypts with it
  - Alice can't decrypt what she's sending to Bob ...
  - ... but then, neither can anyone else (except Bob)

## Public Key / Asymmetric Encryption

- Sender uses receiver's **public** key
  - Advertised to everyone
- Receiver uses complementary **private** key
  - Must be kept secret




4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.21

## Public Key Encryption Details

- Idea:  $K_{\text{public}}$  can be made public, keep  $K_{\text{private}}$  private



- Gives message privacy (restricted receiver):
  - Public keys (secure destination points) can be acquired by anyone/used by anyone
  - Only person with private key can decrypt message
- What about authentication?
  - Use combination of private and public key
  - Alice Bob:  $[(I'm\ Alice)^{A_{\text{private}}}\ \text{Rest of message}]^{B_{\text{public}}}$
  - Provides restricted sender and receiver
- But: how does Alice know that it was Bob who sent her  $B_{\text{public}}$ ? And vice versa...

4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.22

## Public Key Cryptography

- Invented in the 1970s
  - Revolutionized cryptography
  - (Was actually invented earlier by British intelligence)
- How can we construct an encryption/decryption algorithm using a key pair with the public/private properties?
  - Answer: Number Theory
- Most fully developed approach: RSA
  - Rivest / Shamir / Adleman, 1977; RFC 3447
  - Based on modular multiplication of very large integers
  - Very widely used (e.g., ssh, SSL/TLS for https)
- Also mature approach: Elliptic Curve Cryptography (ECC)
  - Based on curves in a Galois-field space
  - Shorter keys and signatures than RSA

4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.23

## Properties of RSA

- Requires generating large, random prime numbers
  - Algorithms exist for quickly finding these (probabilistic!)
- Requires exponentiating very large numbers
  - Again, fairly fast algorithms exist
- Overall, much slower than symmetric key crypto
  - One general strategy: use public key crypto to exchange a (short) symmetric session key
    - » Use that key then with AES or such
- How difficult is recovering  $d$ , the private key?
  - Equivalent to finding prime factors of a large number
    - » Many have tried - believed to be very hard (= brute force only)
    - » (Though quantum computers could do so in polynomial time!)


4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.24

## Simple Public Key Authentication

- Each side need only to know the other side's public key
  - No secret key need be shared
- A encrypts a nonce (random num.)  $x$ 
  - Avoid **replay attacks**, e.g., attacker impersonating client or server
- B proves it can recover  $x$ , generates second nonce  $y$
- A can authenticate itself to B in the same way
- A and B have shared private secrets on which to build private key!
  - We just did secure key distribution!
- Many more details to make this work securely in practice!

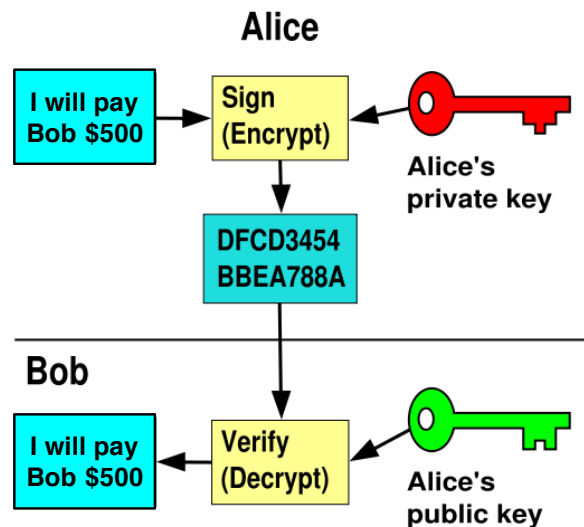


Notation:  $E(m, k)$  – encrypt message  $m$  with key  $k$

## Non-Repudiation: RSA Crypto & Signatures

- Suppose Alice has published public key  $KE$
- If she wishes to prove who she is, she can send a message  $x$  encrypted with her private key  $KD$  (i.e., she sends  $E(x, KD)$ )
  - Anyone knowing Alice's public key  $KE$  can recover  $x$ , verify that Alice must have sent the message
    - It provides a **signature**
  - Alice can't deny it: **non-repudiation**
- Could simply encrypt a hash of the data to sign a document that you wanted to be in clear text
- Note that either of these signature techniques work perfectly well with any data (not just messages)
  - Could sign every datum in a database, for instance

## RSA Crypto & Signatures (cont'd)



## Digital Certificates

- How do you know  $K_E$  is Alice's public key?
- Trusted authority (e.g., Verisign) signs binding between Alice and  $K_E$  with its private key  $KV_{private}$ 
  - $C = E(\{Alice, K_E\}, KV_{private})$
  - $C$ : digital certificate
- Alice: distribute her digital certificate,  $C$
- Anyone: use trusted authority's  $KV_{public}$  to extract Alice's public key from  $C$ 
  - $D(C, KV_{public}) =$   
 $D(E(\{Alice, K_E\}, KV_{private}), KV_{public}) = \{Alice, K_E\}$

## Summary of Our Crypto Toolkit

- If we can securely distribute a key, then
  - Symmetric ciphers (e.g., AES) offer fast, presumably strong confidentiality
- Public key cryptography does away with (potentially major) problem of secure key distribution
  - But: not as computationally efficient
    - » Often addressed by using public key crypto to exchange a **session key**
- Digital signature binds the public key to an entity

4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.29

## Putting It All Together - HTTPS

- What happens when you click on <https://www.amazon.com?>
- `https` = “Use HTTP over SSL/TLS”
  - SSL = Secure Socket Layer
  - TLS = Transport Layer Security
    - » Successor to SSL
  - Provides security layer (authentication, encryption) on top of TCP
    - » Fairly transparent to applications

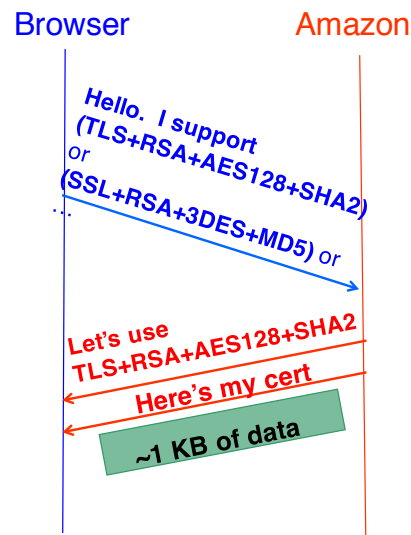
4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.30

## HTTPS Connection (SSL/TLS) (cont'd)

- Browser (client) connects via TCP to Amazon's HTTPS server
- Client sends over list of crypto protocols it supports
- Server picks protocols to use for this session
- Server sends over its certificate
- (all of this is in the clear)



4/27/16

Joseph CSI 62 ©UCB Spring 2016

Lec 24.31

## Inside the Server's Certificate

- Name associated with cert (e.g., Amazon)
- Amazon's **RSA** public key
- A bunch of auxiliary info (physical address, type of cert, expiration time)
- Name of certificate's signatory (who signed it)
- A public-key signature of a hash (**SHA-256**) of all this
  - Constructed using the signatory's private RSA key, i.e.,
    - $\text{Cert} = E_{\text{SHA256}}(K_{\text{A\_public}}, \text{www.amazon.com}, \dots), K_{\text{S\_private}})$
    - »  $K_{\text{A\_public}}$ : Amazon's public key
    - »  $K_{\text{S\_private}}$ : signatory (certificate authority) private key
- ...

4/27/16

Joseph CSI 62 ©UCB Spring 2016

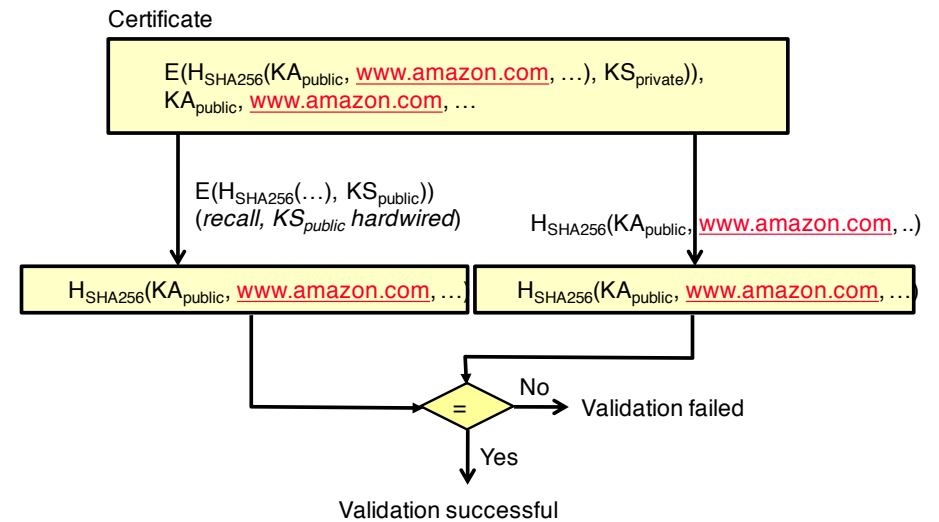
Lec 24.32



## Validating Amazon's Identity

- How does the browser authenticate certificate signatory?
  - Certificates of several certificate authorities (e.g., Verisign) are **hardwired into the browser (or OS)**
- If can't find cert, warn user that site has not been verified
  - And may ask whether to continue
  - Note, can still proceed, just **without authentication**
- Browser uses public key in signatory's cert to decrypt signature
  - Compares with its own **SHA-256** hash of Amazon's cert
- Assuming signature matches, now have high confidence it's indeed Amazon ...
  - ... **assuming signatory is trustworthy**
  - DigiNotar CA breach (July-Sept 2011): Google, Yahoo!, Mozilla, Tor project, Wordpress, ... (**531 total certificates**)

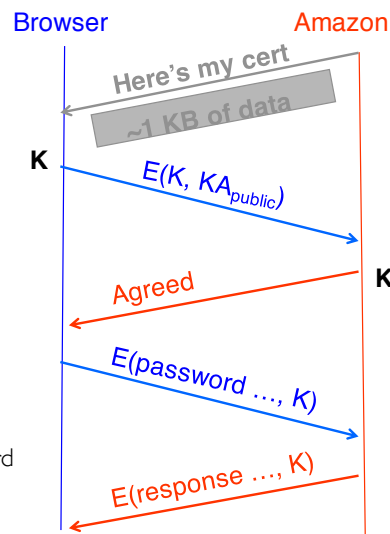
## Certificate Validation



Can also validate using peer approach: <https://www.eff.org/observatory>

## HTTPS Connection (SSL/TLS) cont'd

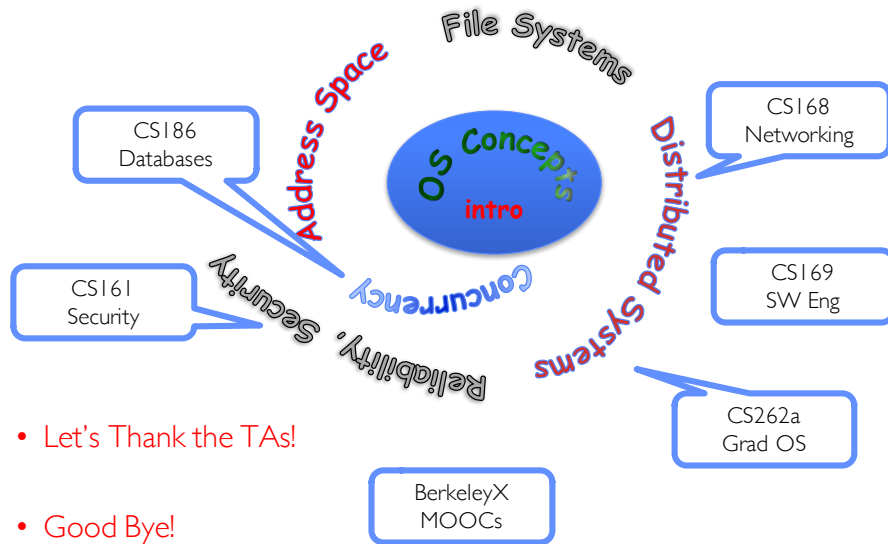
- Browser constructs a random **session key**  $K$  used for data communication
  - Private key for bulk crypto
- Browser encrypts  $K$  using Amazon's public key
- Browser sends  $E(K, KA_{public})$  to server
- Browser displays
- All subsequent comm. encrypted w/ symmetric cipher (e.g., **AES128**) using key  $K$ 
  - E.g., client can authenticate using a password



## Security Summary

- Many more challenges to building secure systems and applications
- No fixed-point solutions
- Adversaries constantly change and adapt
- Defenses must also constantly change and adapt
- Take CS 161

Thank you!



- Let's Thank the TAs!
- Good Bye!