

# Section 10: Intro to I/O and File Systems

April 3, 2016

## Contents

<b>1</b>	<b>Warmup</b>	<b>2</b>
1.1	Short questions . . . . .	2
<b>2</b>	<b>Vocabulary</b>	<b>3</b>
<b>3</b>	<b>Problems</b>	<b>4</b>
3.1	Clock Algorithm . . . . .	4
3.2	Disabling Interrupts . . . . .	4
3.3	Disks . . . . .	5
3.4	Memory Mapped Files . . . . .	5

# 1 Warmup

## 1.1 Short questions

1. (True/False) If a particular IO device implements a blocking interface, then you will need multiple threads to have concurrent operations which use that device.

2. (True/False) For IO devices which receive new data very frequently, it is more efficient to interrupt the CPU than to have the CPU poll the device.

3. (True/False) With SSDs, writing data is straightforward and fast, whereas reading data is complex and slow.

4. (True/False) User applications have to deal with the notion of file blocks, whereas operating systems deal with the finer grained notion of disk sectors.

5. How does the OS catch a null pointer exception ? Trace every action that happens in the OS based on what you have learned in cs162 so far.

## 2 Vocabulary

- **Second Chance Algorithm** A modified form of FIFO that is used to approximate LRU. Each page also has a reference of a use bit to keep track of whether that page has been accessed. It works by looking at the front of the queue as FIFO does, but instead of immediately paging out that page, it checks to see if its referenced bit is set. If it is not set, the page is swapped out. Otherwise, the referenced bit is cleared, the page is inserted at the back of the queue and the process is repeated until a page is swapped out.
- **Clock Algorithm** Clock is a more efficient version of FIFO than Second-chance because pages don't have to be constantly pushed to the back of the list. The clock algorithm keeps a circular list of pages in memory, with the "hand" pointing to the last examined page in the list. When a page fault occurs and no empty frames exist, then the use bit is inspected at the hand's location. If U is 0, the new page is put in place of the page the "hand" points to, otherwise the U bit is cleared. Then, the clock hand is incremented and the process is repeated until a page is replaced.
- **I/O** In the context of operating systems, input/output (I/O) consists of the processes by which the operating system receives and transmits data to connected devices.
- **Controller** The operating system performs the actual I/O operations by communicating with a device controller, which contains addressable memory and registers for communicating with the CPU, and an interface for communicating with the underlying hardware. Communication may be done via programmed I/O, transferring data through registers, or Direct Memory Access, which allows the controller to write directly to memory.
- **Interrupt** One method of notifying the operating system of a pending I/O operation is to send an interrupt, causing an interrupt handler for that event to be run. This requires a lot of overhead, but is suitable for handling sporadic, infrequent events.
- **Polling** Another method of notifying the operating system of a pending I/O operation is simply to have the operating system check regularly if there are any input events. This requires less overhead, and is suitable for regular events, such as mouse input.
- **Response Time** Response time measures the time between a requested I/O operation and its completion, and is an important metric for determining the performance of an I/O device.
- **Throughput** Another important metric is throughput, which measures the rate at which operations are performed over time.
- **Asynchronous I/O** For I/O operations, we can have the requesting process sleep until the operation is complete, or have the call return immediately and have the process continue execution and later notify the process when the operation is complete.
- **Memory-Mapped File** A memory-mapped file is a segment of virtual memory which has been assigned a direct byte-for-byte correlation with some portion of a file or file-like resource. This resource is typically a file that is physically present on-disk, but can also be a device, shared memory object, or other resource that the operating system can reference through a file descriptor. Once present, this correlation between the file and the memory space permits applications to treat the mapped portion as if it were primary memory.
- **Memory-Mapped I/O** Memory-mapped I/O (not to be confused with memory-mapped file I/O) uses the same address bus to address both memory and I/O devices the memory and registers of the I/O devices are mapped to (associated with) address values. So when an address is accessed by the CPU, it may refer to a portion of physical RAM, but it can also refer to memory of the I/O device. Thus, the CPU instructions used to access the memory can also be used for accessing devices.

### 3 Problems

#### 3.1 Clock Algorithm

Suppose that we have a 32-bit virtual address split as follows:

10 Bits	10 Bits	12 Bits
Table ID	Page ID	Offset

Show the format of a PTE complete with bits required to support the clock algorithm.

For this problem, assume that physical memory can hold at most four pages. What pages remain in memory at the end of the following sequence of page table operations and what are the use bits set to for each of these pages:

- Page A is paged in
- Page B is paged in
- Page C is paged in
- Page A is accessed
- Page C is accessed
- Page D is paged in
- Page B is accessed
- Page D is accessed
- Page A is accessed
- Page E is paged in
- Page F is paged in

#### 3.2 Disabling Interrupts

We looked at disabling CPU interrupts as a simple way to create a critical section in the kernel. Name a drawback of this approach when it comes to I/O devices.

### 3.3 Disks

What are the major components of disk latency? Explain each one.

In class we said that the operating system deals with bad or corrupted sectors. Some disk controllers magically hide failing sectors and re-map to back-up locations on disk when a sector fails.

If you had to choose where to lay out these back-up sectors on disk - where would you put them? Why?

How do you think that the disk controller can check whether a sector has gone bad?

Can you think of any drawbacks of hiding errors like this from the operating system?

### 3.4 Memory Mapped Files

Memory-mapped files, either as the fundamental way to access files or as an additional feature, is supported by most OS-s. Consider the Unix mmap system call that has the form (somewhat simplified):

```
void * mmap(void * addr, size_t len, int prot, int flags, int fildes, off_t offset);
```

where data is being mapped from the currently open file *fildes* starting at the position in the file described by *offset*; *len* is the size of the part of the file that is being mapped into the process address space; *prot* describes whether the mapped part of the is readable or writable. The return value, *addr*, is the address in the process address space of the mapped file region.

When the process subsequently references a memory location that has been mapped from the file for the first time, what operations happen in the operating system? What happens upon subsequent accesses?

Assume that you have a program that will read sequentially through a very large file, computing some summary operation on all the bytes in the file. Compare the efficiency of performing this task when you are using conventional read system calls versus using `mmap`.

Assume that you have a program that will read randomly from a very large file. Compare the efficiency of performing this task when you are using conventional read and `lseek` system calls versus using `mmap`.